

# Deep Learning lecture 4 Energy-Based Model

Yi Wu, IIIS

Spring 2025

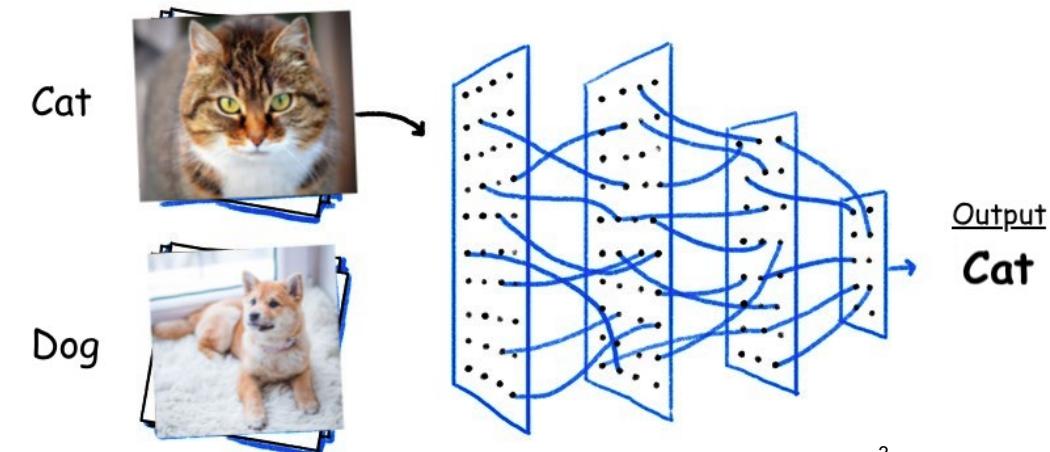
Mar-10

# Logistics

- Coding Project 2 due in 1 week
  - Use local compute for coding & Colab for testing
  - Cloud for long-term training
  - Any questions can be posted in Dingding channel
  - Be aware of your model size and computation (flops)!
  - Check out those famous models and works!

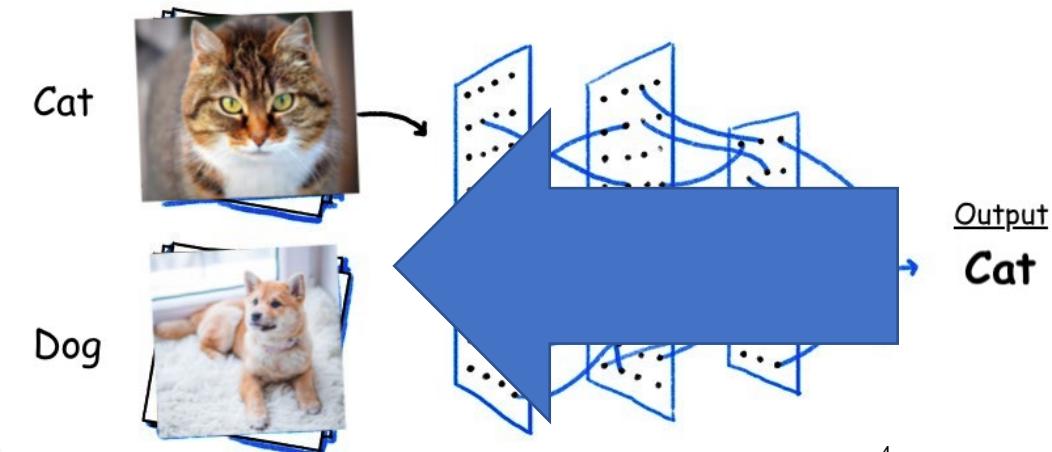
# Story So Far

- History
  - Lecture 1
    - first neural network (1943) to recent advances in deep learning
- Supervised Learning (Classification)
  - Lecture 2
    - MLP and basic components; Backpropagation
  - Lecture 3
    - Algorithms, Tricks and Architecture
- Discriminative Model
  - $P(y|X)$
  - Labeled data;  $X \rightarrow y$



# Afterwards

- What if we want to generate  $X$ ?
  - E.g., Ask the neural network to generate a cat!
- Generative Model
  - $P(X, y) = P(y) * P(X|y)$
  - Or just  $P(X)$
- Lecture 4~7
  - Deep Generative Models
  - Different approaches to model  $P(X)$



# Today's Lecture: Energy-Based Models

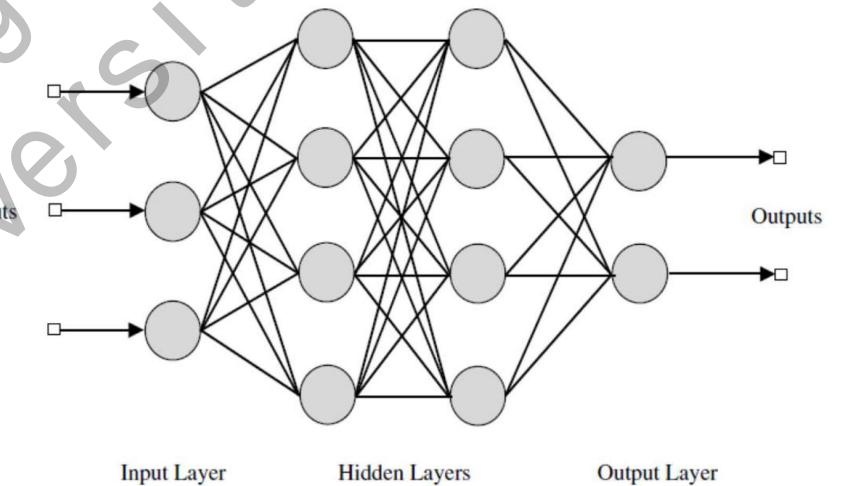
- A particularly flexible and general form of ***generative model***
- Part 1: Hopfield Network
  - The simplest model that can memorize and generate patterns
- Part 2: Boltzmann Machine
  - The first deep generative model
- Part 3: General Energy-Based Models & Sampling Methods

# Today's Lecture: Energy-Based Models

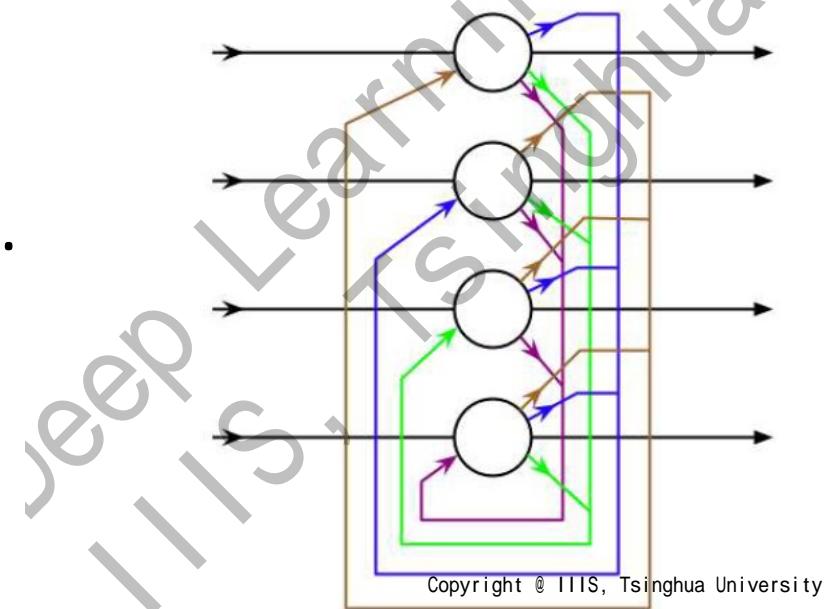
- A particularly flexible and general form of ***generative model***
- Part 1: Hopfield Network
  - The simplest model that can memorize and generate patterns
- Part 2: Boltzmann Machine
  - The first deep generative model
- Part 3: General Energy-Based Models & Sampling Methods

# Classification

- Recap: Classification
  - Layer-by-layer computation
  - Computation Graph: Directed Acyclic Graph
  - Feedforward networks



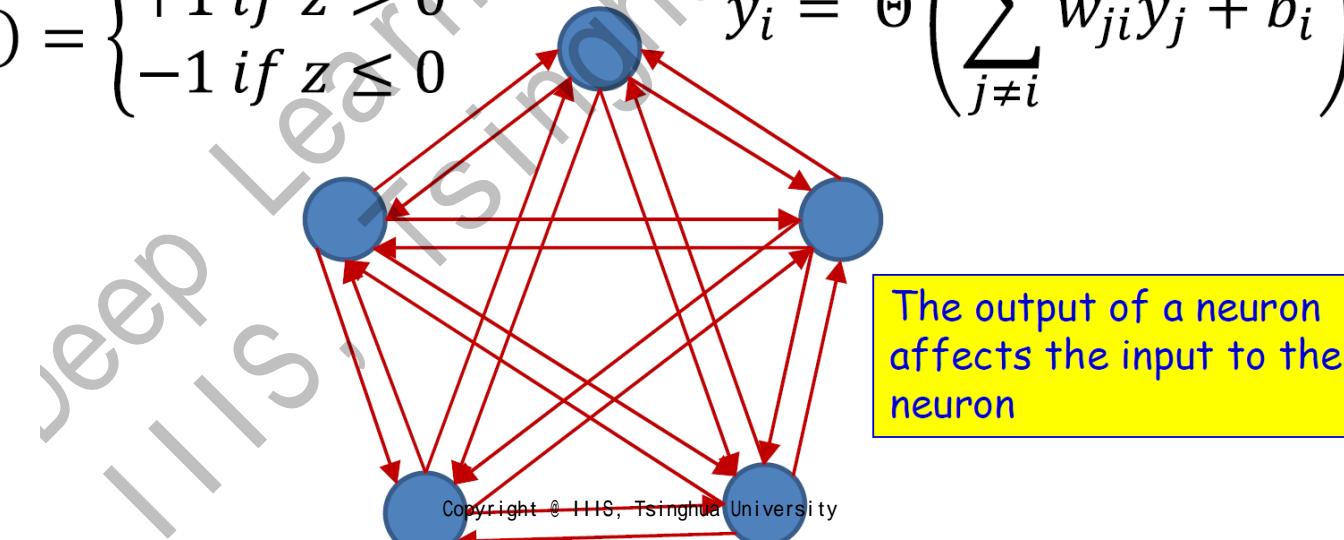
- What about ...
  - Loops!



# A Loopy Network

- A “fully-connected” network
  - Each neuron receives inputs from all the other neurons
  - $y_i = +1 \text{ or } -1$  with hard thresholding

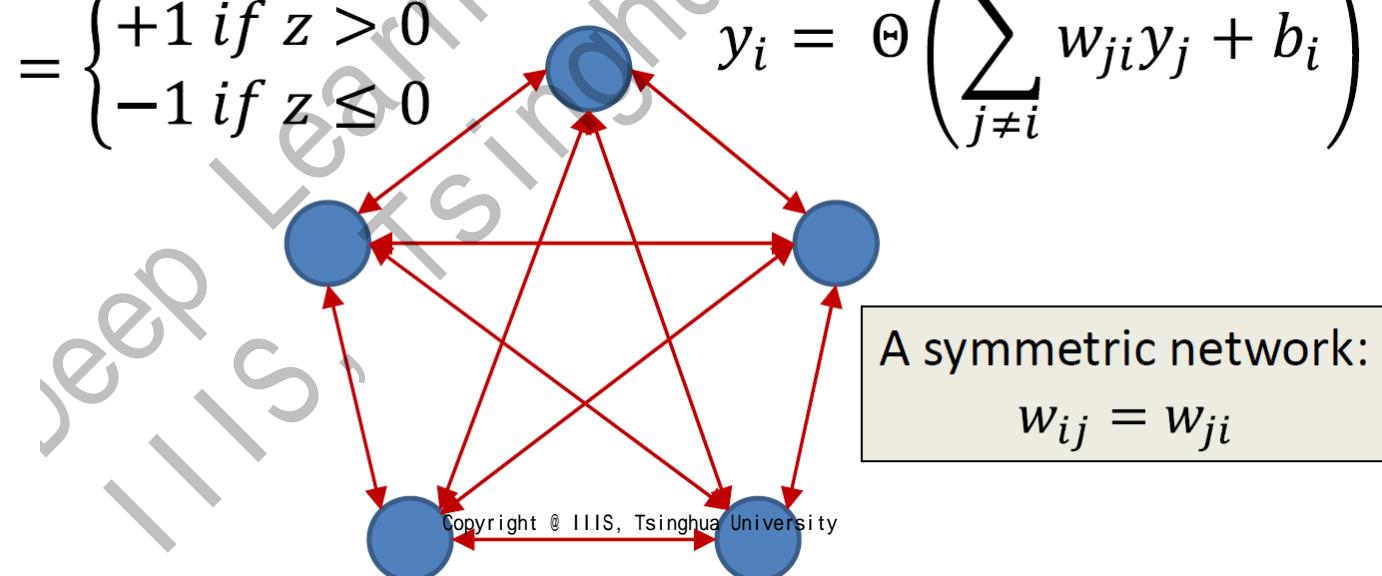
$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$$



# Hopfield Network

- A “fully-connected” network
  - Each neuron receives inputs from all the other neurons
  - $y_i = +1 \text{ or } -1$  with hard thresholding
  - Symmetric weights

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji}y_j + b_i\right)$$

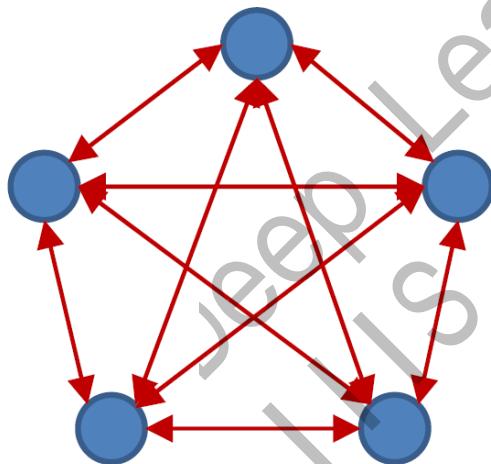


# Hopfield Network

- A Hopfield Network may not be stable!
  - At each time each neuron receives a “field”  $z_i = \sum_{j \neq i} w_{ji}y_j + b_i$
  - If the sign of neuron matches the sign of the field, it flips

$$y_i \leftarrow -y_i \text{ if } y_i \left( \sum_{j \neq i} w_{ji}y_j + b_i \right) < 0$$

- This can further cause other neurons to flip!



$$y_i = \Theta \left( \sum_{j \neq i} w_{ji}y_j + b_i \right)$$

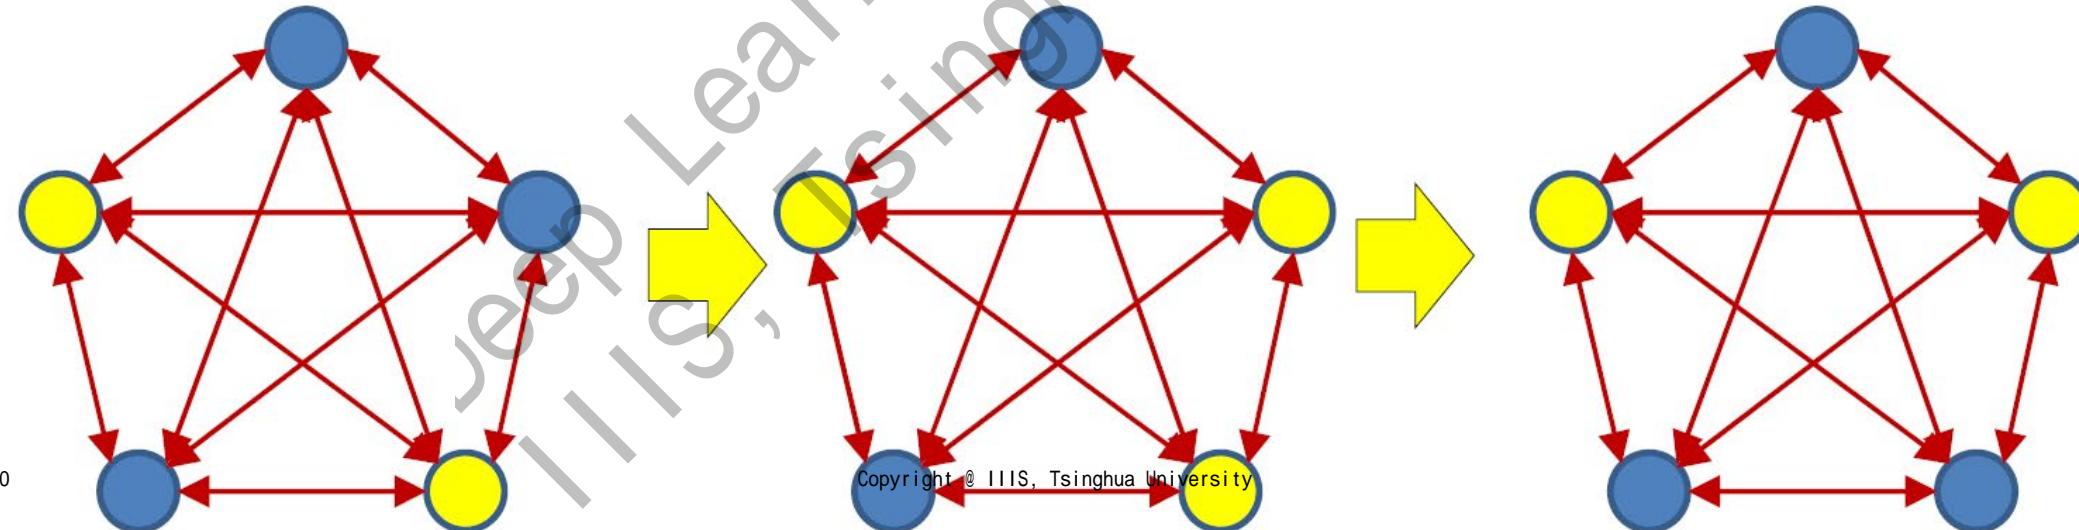
$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

# Hopfield Network

- Neurons flip if its sign does not match its local “field”
  - $y_i \leftarrow -y_i$  if  $y_i(\sum_{j \neq i} w_{ji}y_j + b_i) < 0$  for all neurons
  - Repeat until no neuron flips
  - Will this process converge?

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

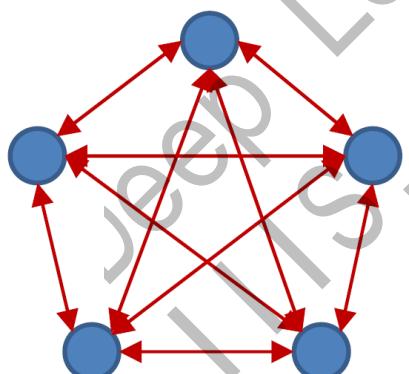
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji}y_j + b_i\right)$$



# Hopfield Network

- Let  $y_i^-$  denote the value of  $y_i$  before a “flip”
- Let  $y_i^+$  denote the value of  $y_i$  after a “flip”
- If  $y_i^- (\sum_{j \neq i} w_{ji} y_j + b_i) \geq 0$ , nothing happen

$$y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) - y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) = 0$$



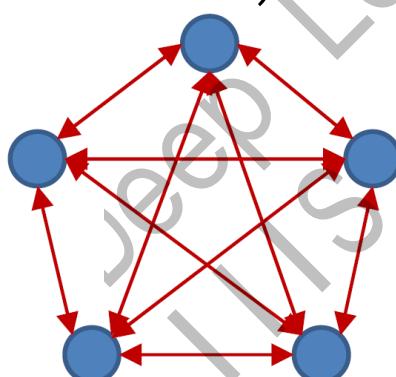
$$y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

Copyright © IIIS, Tsinghua University

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

# Hopfield Network

- Let  $y_i^-$  denote the value of  $y_i$  before a “flip”
  - Let  $y_i^+$  denote the value of  $y_i$  after a “flip”
  - If  $y_i^- (\sum_{j \neq i} w_{ji} y_j + b_i) \geq 0$ , nothing happen
  - If  $y_i^- (\sum_{j \neq i} w_{ji} y_j + b_i) < 0$ ,  $y_i^+ = -y_i^-$
- $$y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) - y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) = 2y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$
- $$y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$



Copyright © IIIS, Tsinghua University

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

# Hopfield Network

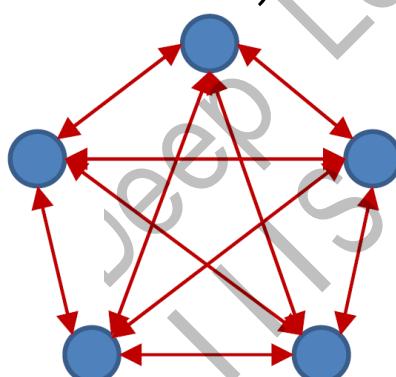
- Let  $y_i^-$  denote the value of  $y_i$  before a “flip”
- Let  $y_i^+$  denote the value of  $y_i$  after a “flip”
- If  $y_i^- (\sum_{j \neq i} w_{ji} y_j + b_i) \geq 0$ , nothing happen

- If  $y_i^- (\sum_{j \neq i} w_{ji} y_j + b_i) < 0$ ,  $y_i^+ = -y_i^-$

$$y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) - y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) = 2y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

**Positive!**

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$



# Hopfield Network

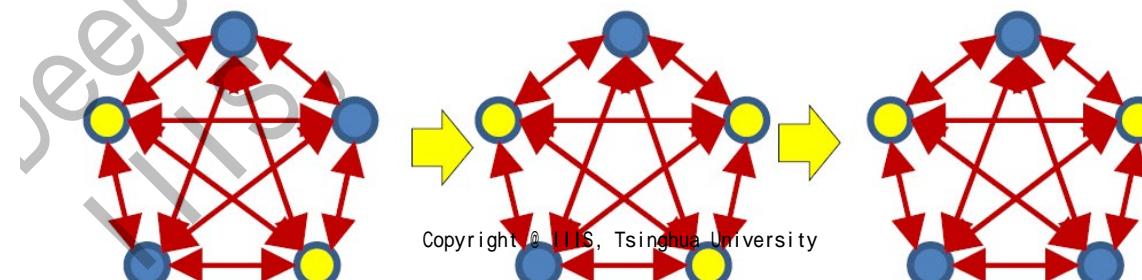
- Consider the sum over every pair of neurons (assume  $w_{ii} = 0$ )

$$D(y_1, \dots, y_N) = \sum_{i < j} y_i w_{ij} y_j + y_i b_i$$

- Any flip that changes  $y_i^-$  to  $y_i^+$  increases  $D(y_1, \dots, y_N)$

$$\Delta D = D(\dots, y_i^+, \dots) - D(\dots, y_i^-, \dots) = 2y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) > 0$$

- Convergence?



# Hopfield Network

- $D$  is upper-bounded (we only change  $y_i$ )

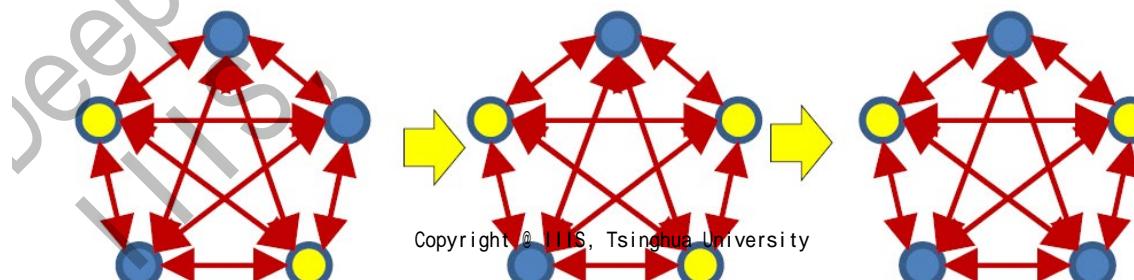
$$D(y_1, \dots, y_N) = \sum_{i < j} w_{ij} y_i y_j + \sum_i b_i y_i \leq \sum_{i < j} |w_{ij}| + \sum_i |b_i|$$

- $\Delta D$  is lower-bounded

$$\Delta D_{\min} = \min_{i, \{y_j\}} 2 \left| \sum_j w_{ij} y_j + b_i \right| > 0$$

- $\{y_i\}$  converges with a finite number of iterations!

- $\{y_i\}$ : state



# Hopfield Network

- The ***Energy*** of Hopfield Network

$$E = -D = - \sum_{i < j} w_{ij} y_i y_j - \sum_i b_i y_i$$

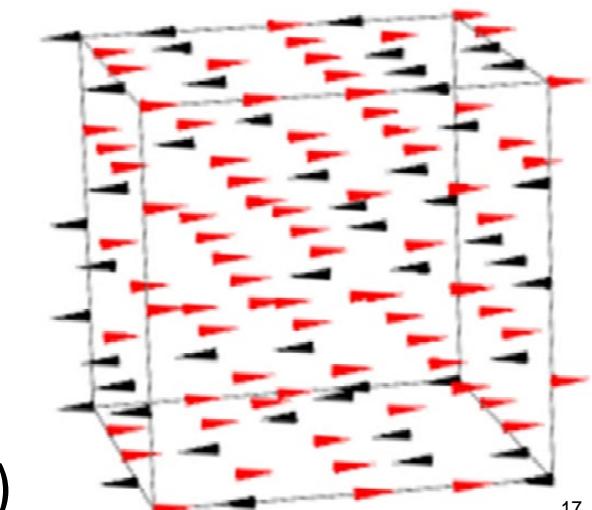
- The evolution of Hopfield network always decreases its energy!

- The concept of ***Energy***

- Magnetic dipoles in a disordered magnetic material
- Each dipole tries to align itself to the local field
- Field at a particular dipole  $f(p_i)$ ,  $p_i$  is the position of  $x_i$

$$f(p_i) = \sum_{j \neq i} J_j x_j + b_i$$

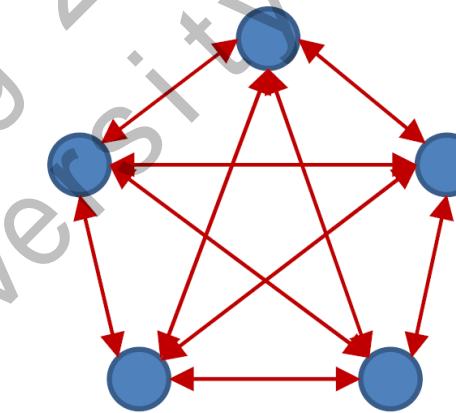
- ***Ising model*** of magnetic materials (Ising and Lenz, 1924)



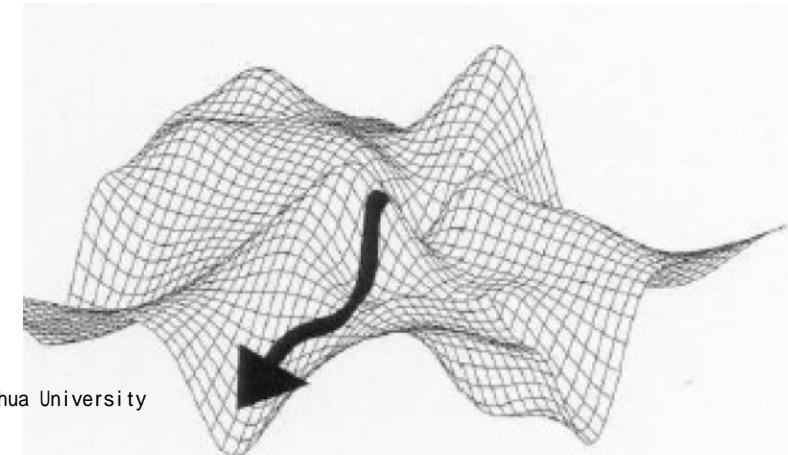
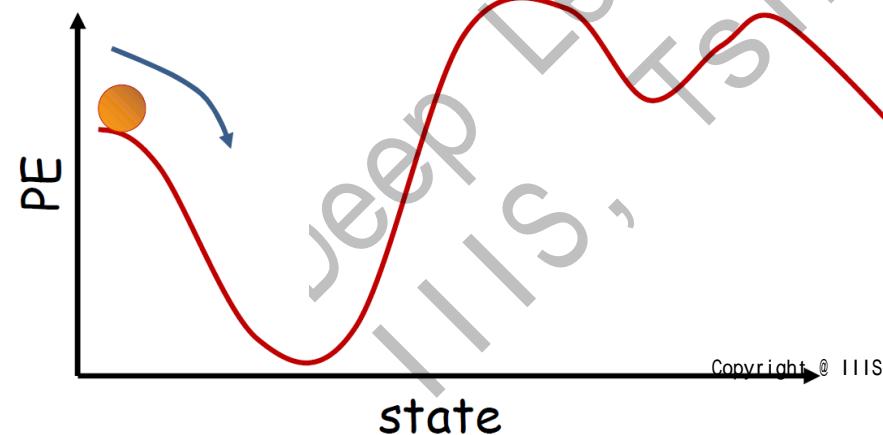
# Hopfield Network: Pattern Generation

- The Hopfield network (simplified)

$$E = - \sum_{i < j} w_{ij} y_i y_j$$



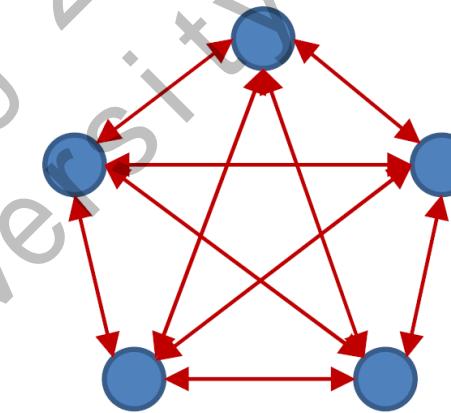
- Network evolution arrives at a local optimum in the energy contour
  - Every change in the network state  $Y$  decreases the energy  $E$
- Any small jitter from this stable state returns it to the stable state***



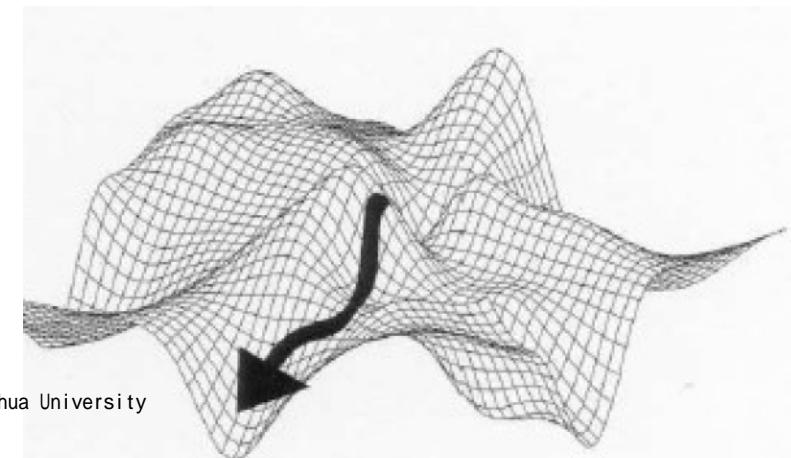
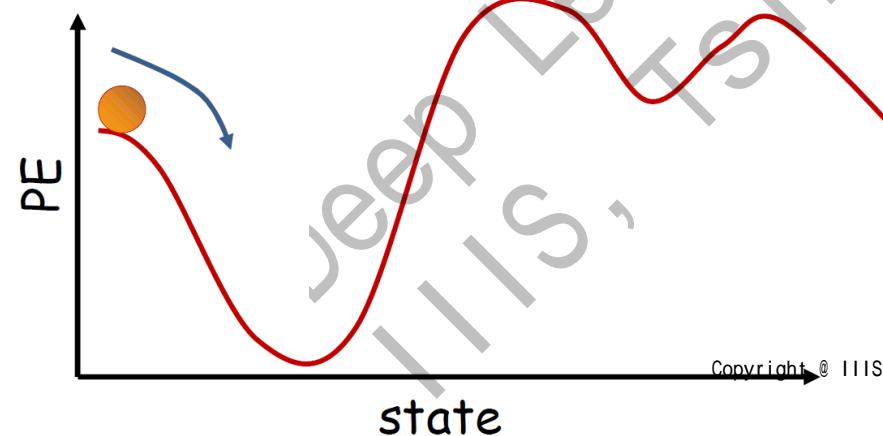
# Hopfield Network: Pattern Generation

- The Hopfield network (simplified)

$$E = - \sum_{i < j} w_{ij} y_i y_j$$

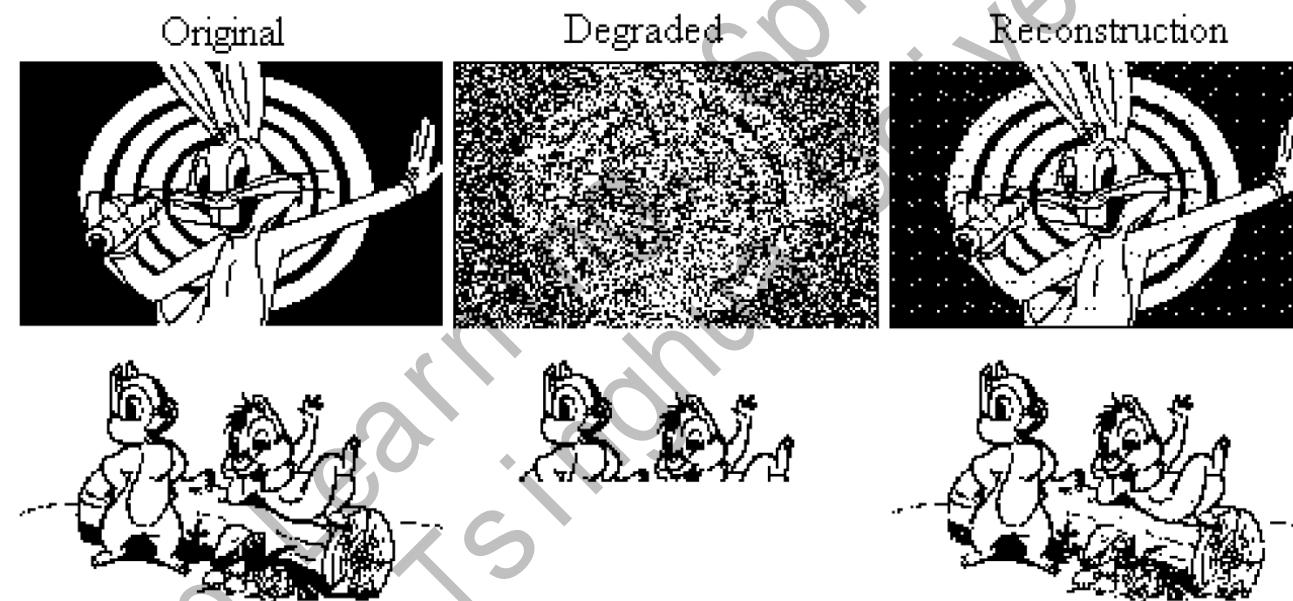


- Each local optimum state is a “stored” pattern
  - If the network is initialized close to a stored pattern, it evolves to the pattern
- Associated Memory (content addressable memory)***



# Hopfield Network: Pattern Generation

- Image Reconstruction by Hopfield Network (1982)



Hopfield network reconstructing degraded images  
from noisy (top) or partial (bottom) cues.

- *How can we store the desired patterns?*

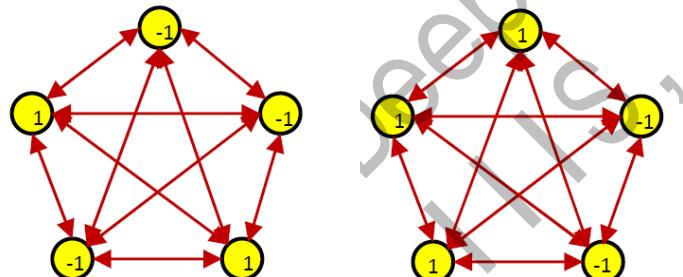
# Hopfield Network: Training

- Let's teach the network to store this image
  - $N$  pixels  $\rightarrow N$  neurons
  - Symmetric weights  $\rightarrow \frac{1}{2}N(N - 1)$  parameters to learn
    - We omit bias terms for simplicity
- Design  $\{w_{ij}\}$  such that the energy is at a local minimum for a desired pattern  $y$ 
  - Hebbian Learning Rule  $w_{ij} \leftarrow y_i y_j$  (1949)
  - $E = -\sum_{i < j} w_{ij} y_i y_j = -\frac{1}{2}N(N - 1) \rightarrow$  lowest possible energy!



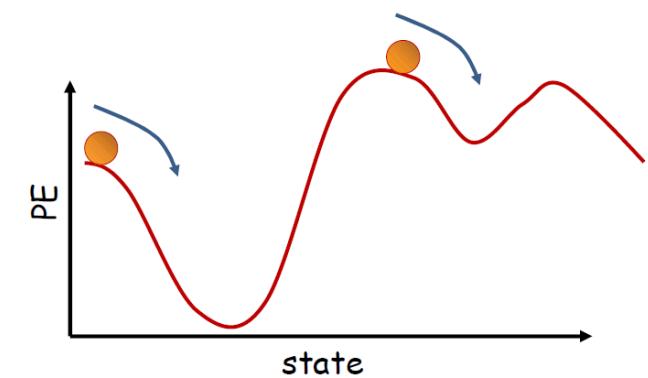
# Hopfield Network: Training

- Let's teach the network to store this image
  - $N$  pixels  $\rightarrow N$  neurons
  - Symmetric weights  $\rightarrow \frac{1}{2}N(N - 1)$  parameters to learn
    - We omit bias terms for simplicity
- Design  $\{w_{ij}\}$  such that the energy is at a local minimum for a desired pattern  $y$ 
  - **Redundancy!**  $y$  &  $-y$  will be both stored



$$E = - \sum_i \sum_{j < i} w_{ji} y_j y_i$$

Copyright © IIIS, Tsinghua University



# Hopfield Network: Training

- What if we want to store ***multiple*** patterns?

- $P = \{y^p\}$   $N_p$  patterns
- Hebbian Learning Rule

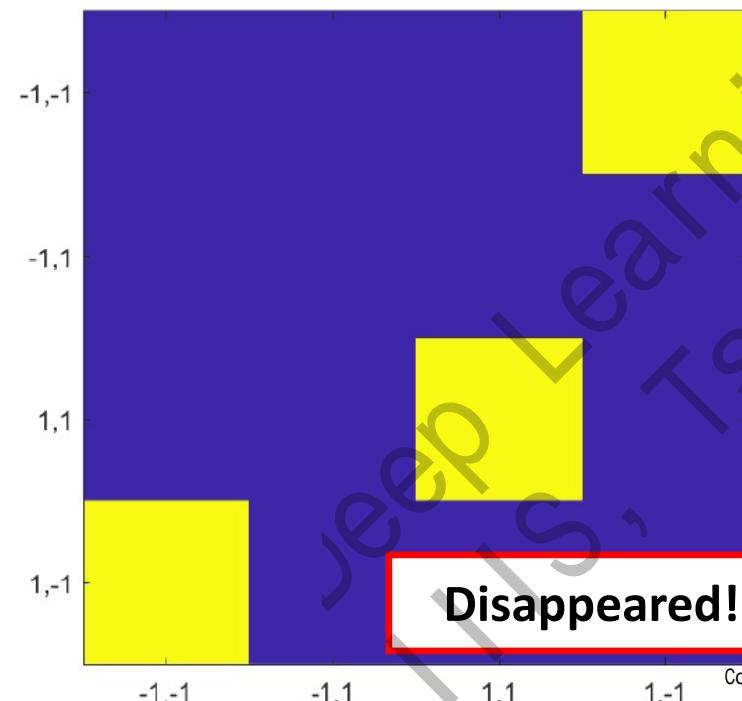
$$w_{ij} = \frac{1}{N_p} \sum_p y_i^p y_j^p$$

- The issue of Hebbian Learning
  - Spurious local optima

# Hopfield Network: Training

- Example: 4-dimensional Hopfield Network with Hebbian Learning
  - Three patterns to store
    - *Let's assume the value of each neuron is 1 or -1*

*Left:  
desired  
patterns*



Copyright © IIIS, Tsinghua University

Copyright © IIIS, Tsinghua University

Copyright © IIIS, Tsinghua University

"Fake" memory

3 nonorthogonal patterns

*Right:  
stored  
patterns*

# Hopfield Network: Training

- We want to construct a network with desired **stable** local optimum
  - A pattern can be recovered after 1-bit change
- For a specific set of  $K$  patterns, we can always build a network for which all patterns are stable provided  $K \leq N$ 
  - Mostafa and St. Jacques (1985)
  - For large  $N$ , the upper bound on  $K$  is actually  $\frac{N}{4 \log N}$
  - McEliece et. al. (1987)
  - Still possible with undesired local minimum
- **How can we find the weights?**
  - $K$  patterns to be stored
  - Avoid undesired local minimum as much as we can

# Hopfield Network: Optimization

- Problem Formulation
  - Desired patterns  $P = \{y^p\}$
  - Energy function  $E(y) = -\frac{1}{2}y^T W y$  (we omit bias term for simplicity)
- Objective for  $W$ 
  - Minimize  $E$  for all  $y^p$
  - It should also maximize  $E$  for all non-desired patterns!

$$W = \arg \min_W \sum_{y \in P} E(y) - \sum_{y' \notin P} E(y')$$

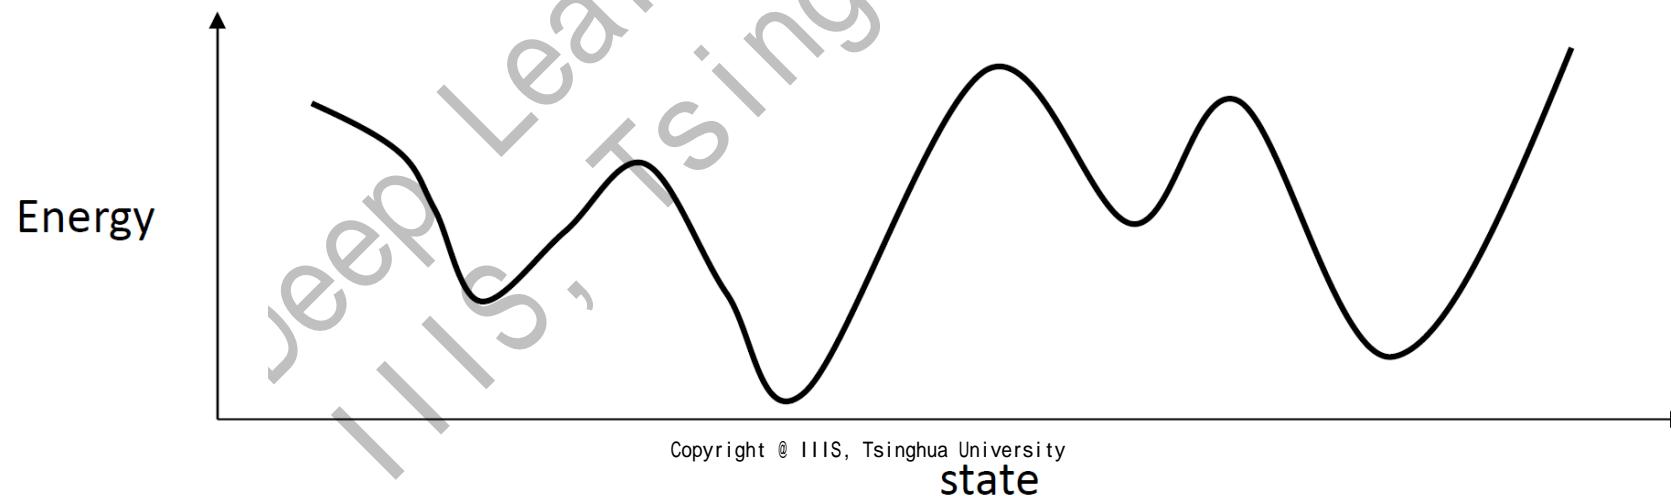
- Gradient Descent

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P} y'y'^T \right)$$

# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P} y'y'^T \right)$$

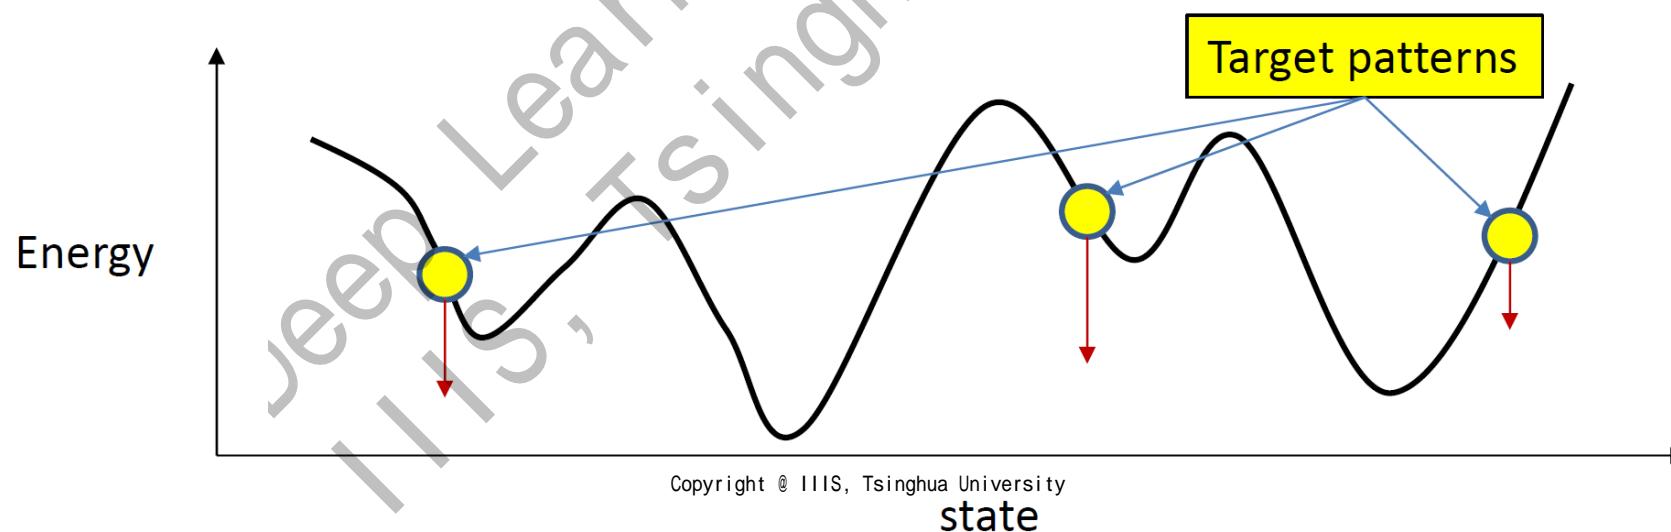


# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P} y'y'^T \right)$$

- The first term is minimizing the energy of desired patterns!

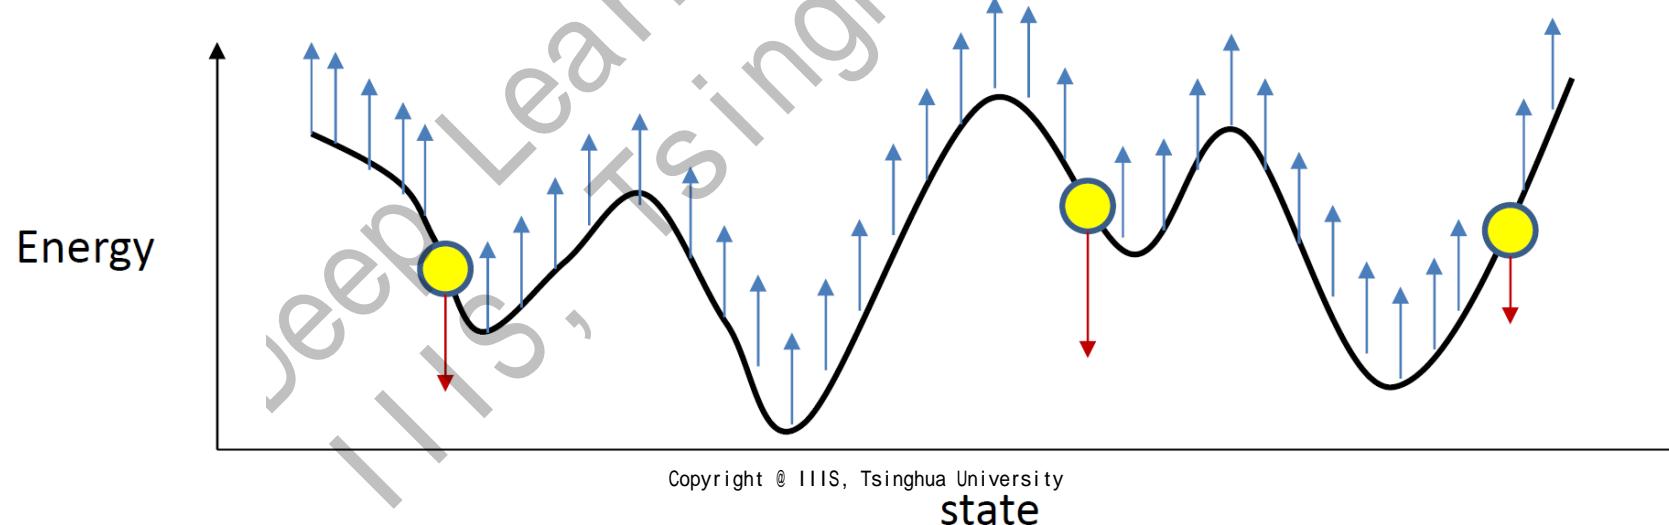


# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P} y'y'^T \right)$$

- The second term essentially raises all the patterns in the space
  - Issue??

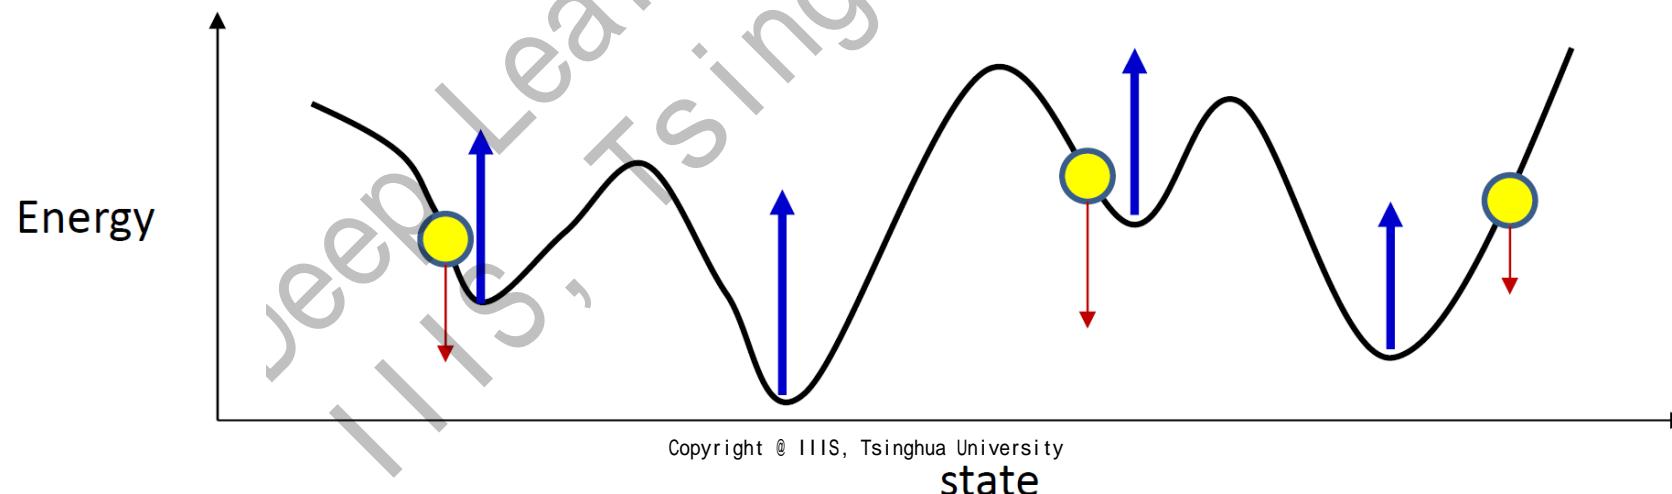


# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ & } y' \in \text{Valley}} y'y'^T \right)$$

- Let's just focus on the valleys!

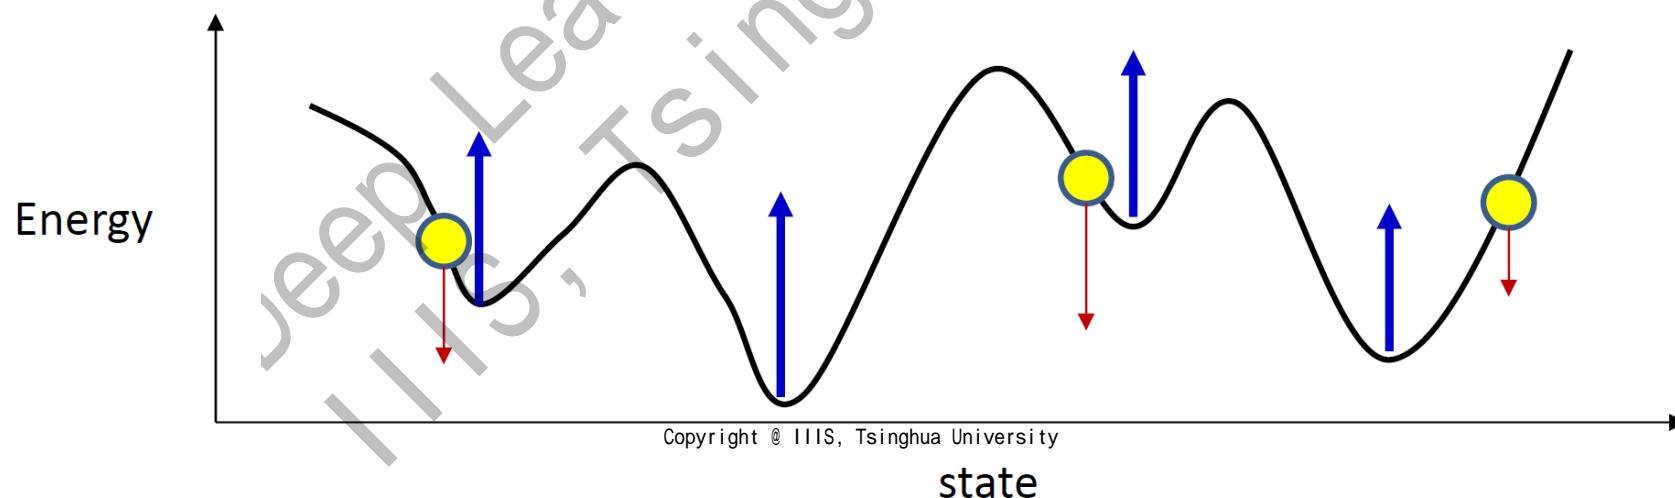


# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ & } y' \in Valley} y'y'^T \right)$$

- Let's just focus on the valleys!
- **But how can we find the valleys?**

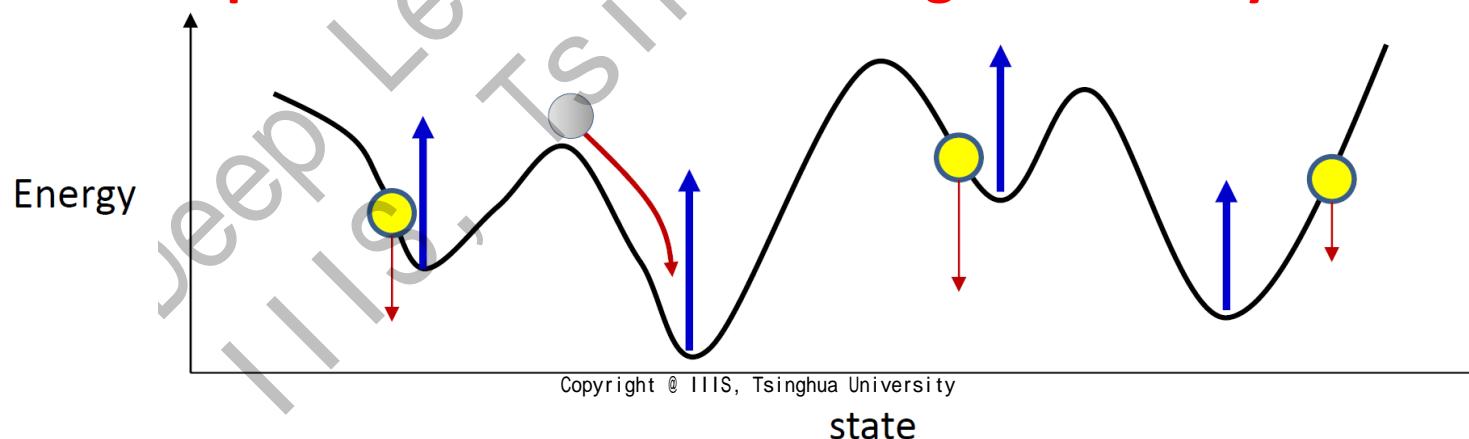


# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ & } y' \in Valley} y'y'^T \right)$$

- Let's just focus on the valleys!
- But how can we find the valleys?
- **Evolution of Hopfield Network will converge to a valley**



# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ & } y' \in Valley} y'y'^T \right)$$

- Compute outer-products of desired patterns  $y$
- Randomly initialize  $y'$  for multiple times
  - Run evolution for random  $y'$  until convergence
  - Calculate outer-product of  $y'$
- Compute gradient and update  $W$

# Hopfield Network: Optimization

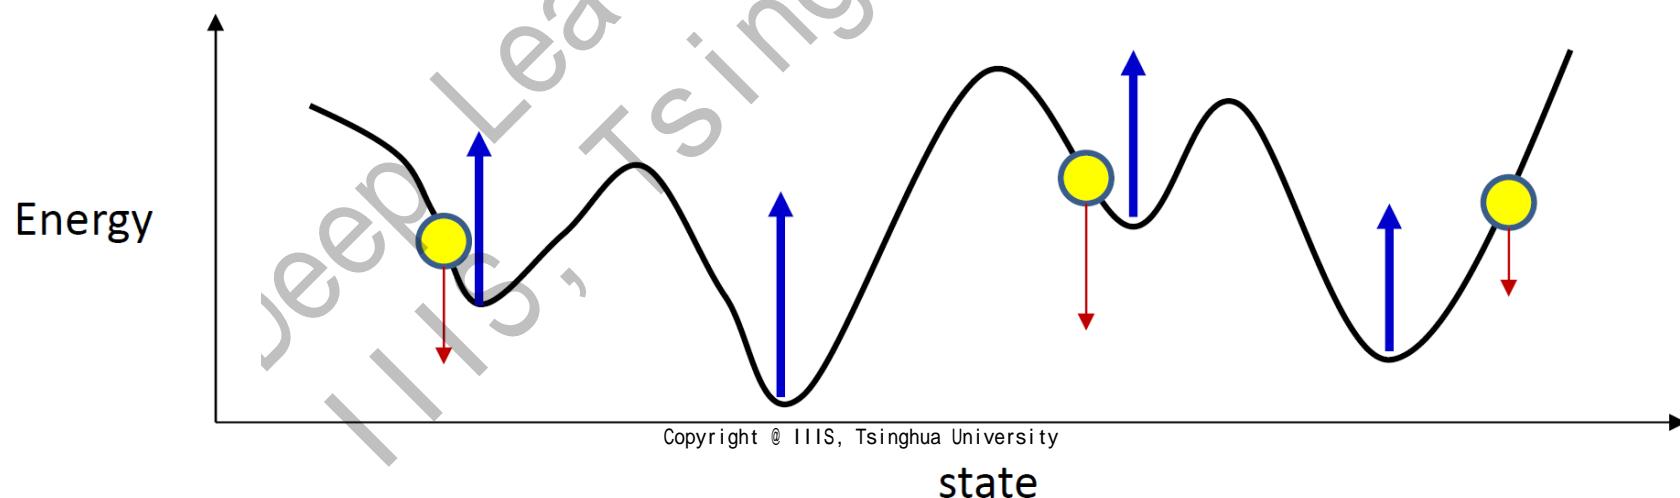
- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ & } y' \in Valley} y'y'^T \right)$$

- Compute outer-products of desired patterns  $y$
- **Randomly initialize**  $y'$  for multiple times
  - Run evolution for random  $y'$  until convergence
  - Calculate outer-product of  $y'$
- Compute gradient and update  $W$
- **Valleys are NOT equivalently important...**

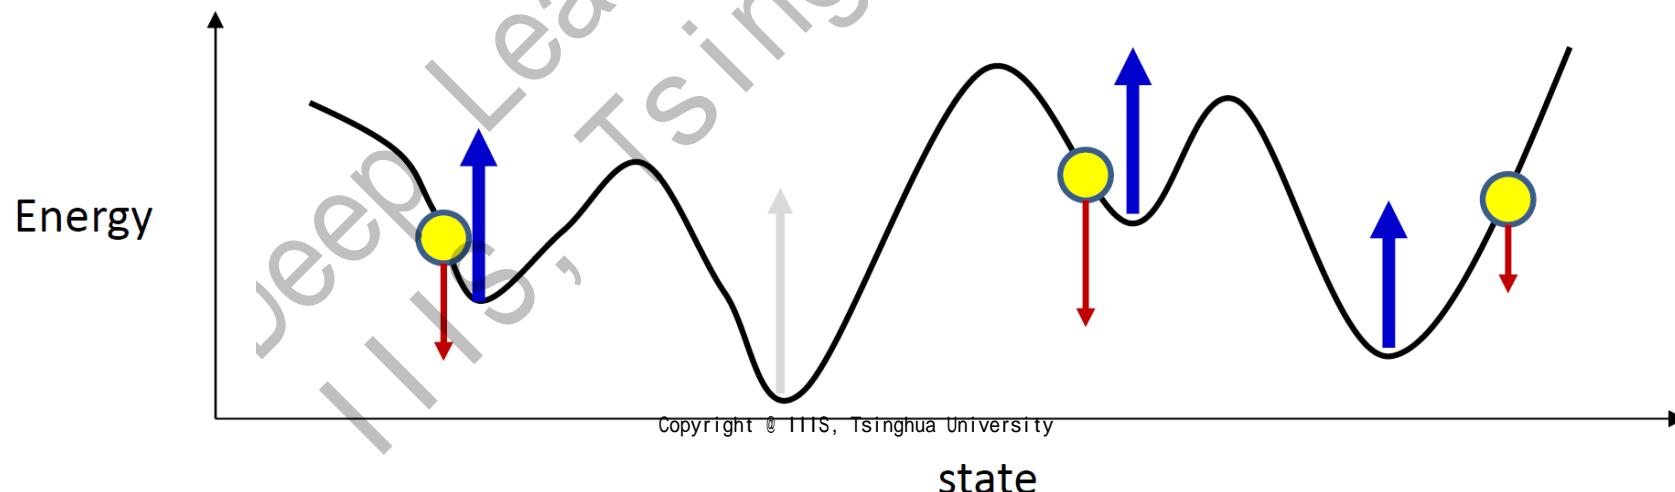
# Hopfield Network: Optimization

- Which valleys are important?
- Primary object: ensure desired patterns stable
  - We want to ensure desired patterns are in broad valleys



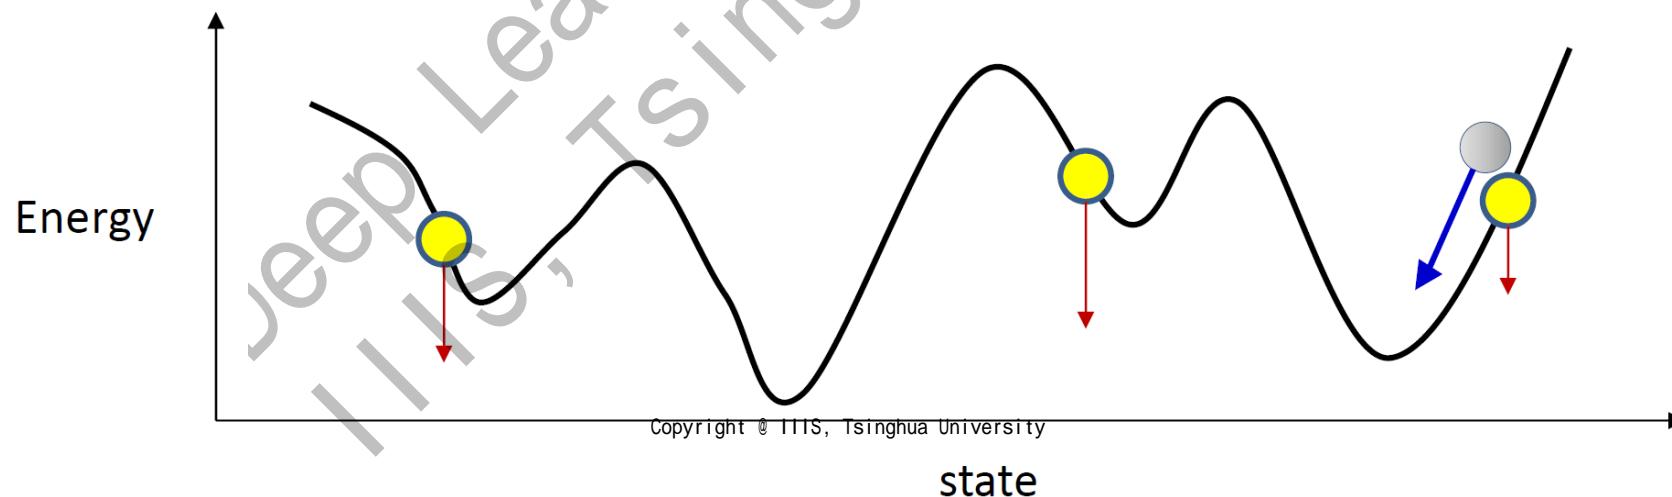
# Hopfield Network: Optimization

- Which valleys are important?
- Primary object: ensure desired patterns stable
  - We want to ensure desired patterns are in broad valleys
  - **Spurious valleys around desired patterns are more important to eliminate**



# Hopfield Network: Optimization

- Which valleys are important?
- Primary object: ensure desired patterns stable
  - We want to ensure desired patterns are in broad valleys
  - Spurious valleys around desired patterns are more important to eliminate
  - **Evolution from desired patterns**



# Hopfield Network: Optimization

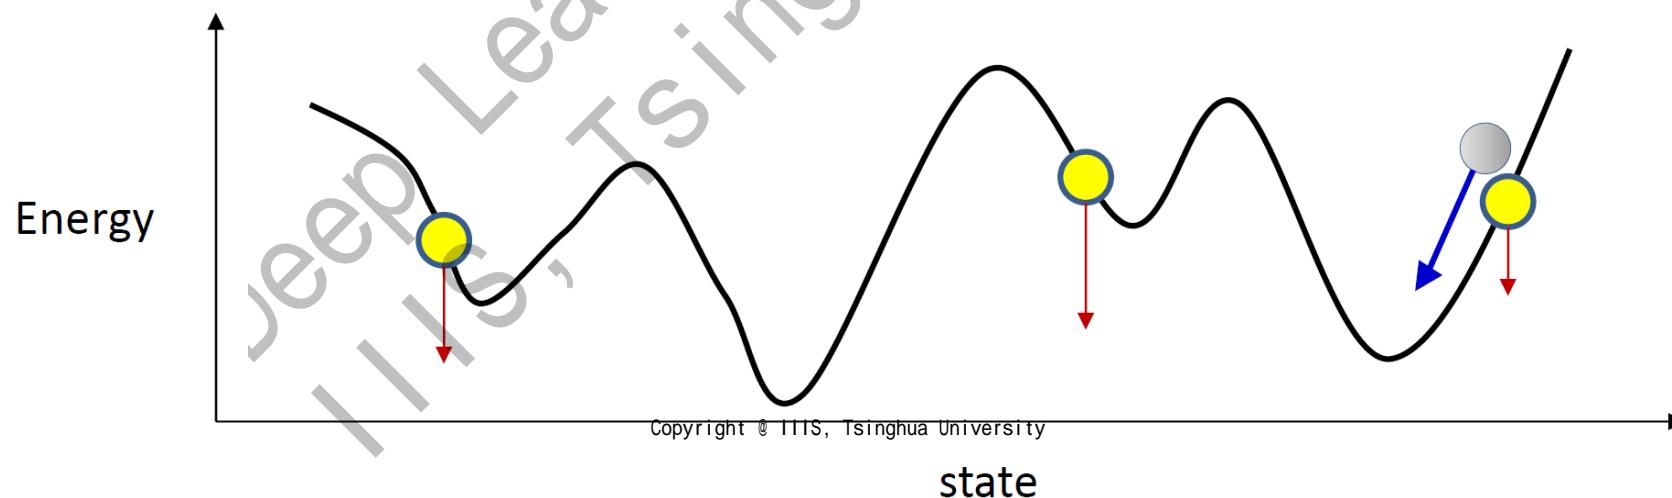
- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ & } y' \in Valley} y'y'^T \right)$$

- Compute outer-products of desired patterns  $y$
- Initialize  $y'$  by all the desired patterns
  - Run evolution for random  $y'$  until convergence
  - Calculate outer-product of  $y'$
- Compute gradient and update  $W$
- Still issues?

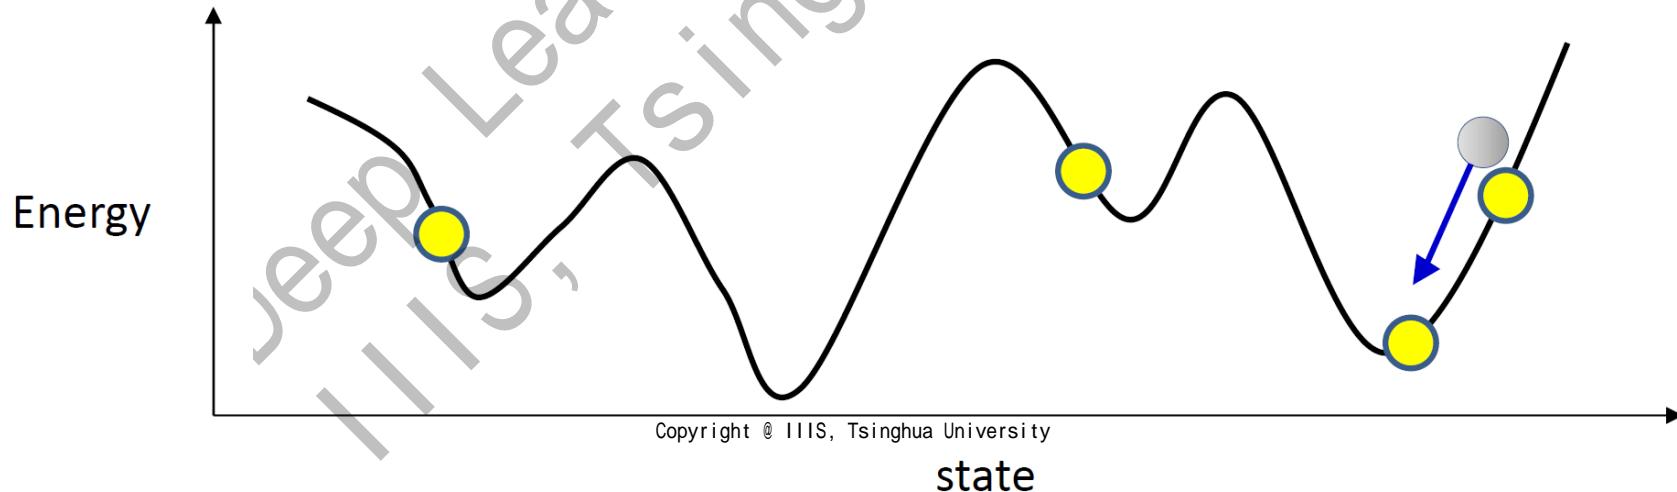
# Hopfield Network: Optimization

- Recap: we raise the valleys next to the desired patterns



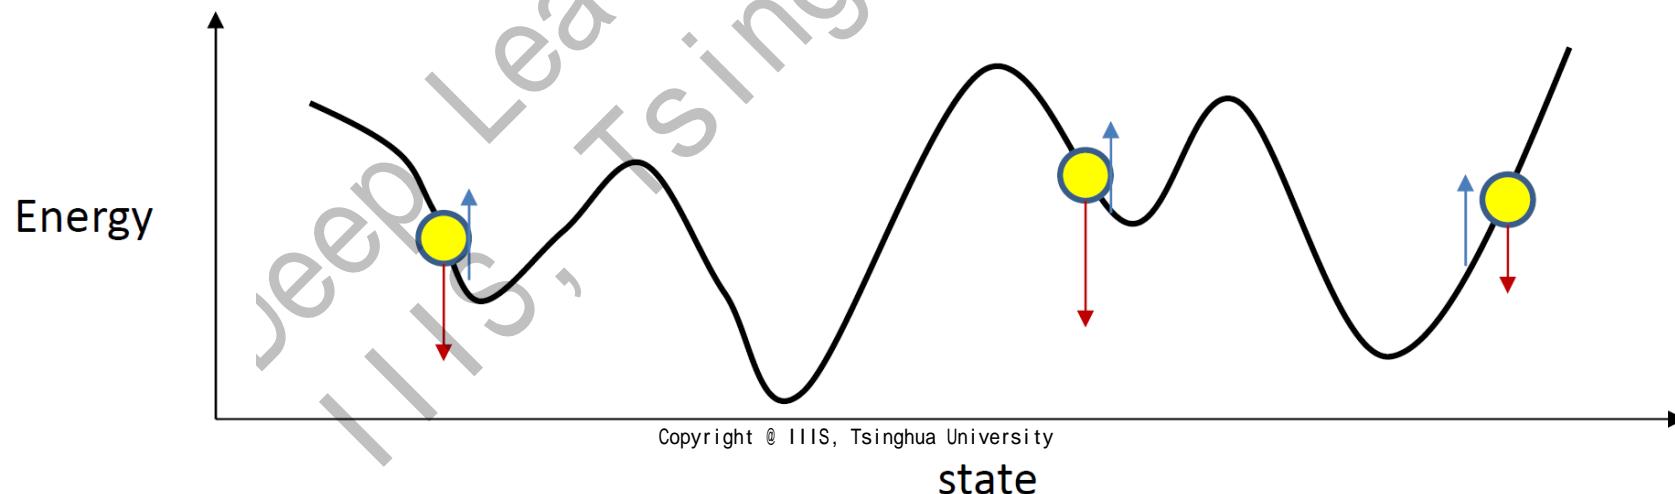
# Hopfield Network: Optimization

- Recap: we raise the valleys next to the desired patterns
- What if a pattern is close to the valley?
  - Naively forcing a valley to raise may hurt the learned representation
  - Particularly challenging when  $y$  are continuously valued (e.g., tanh activation)



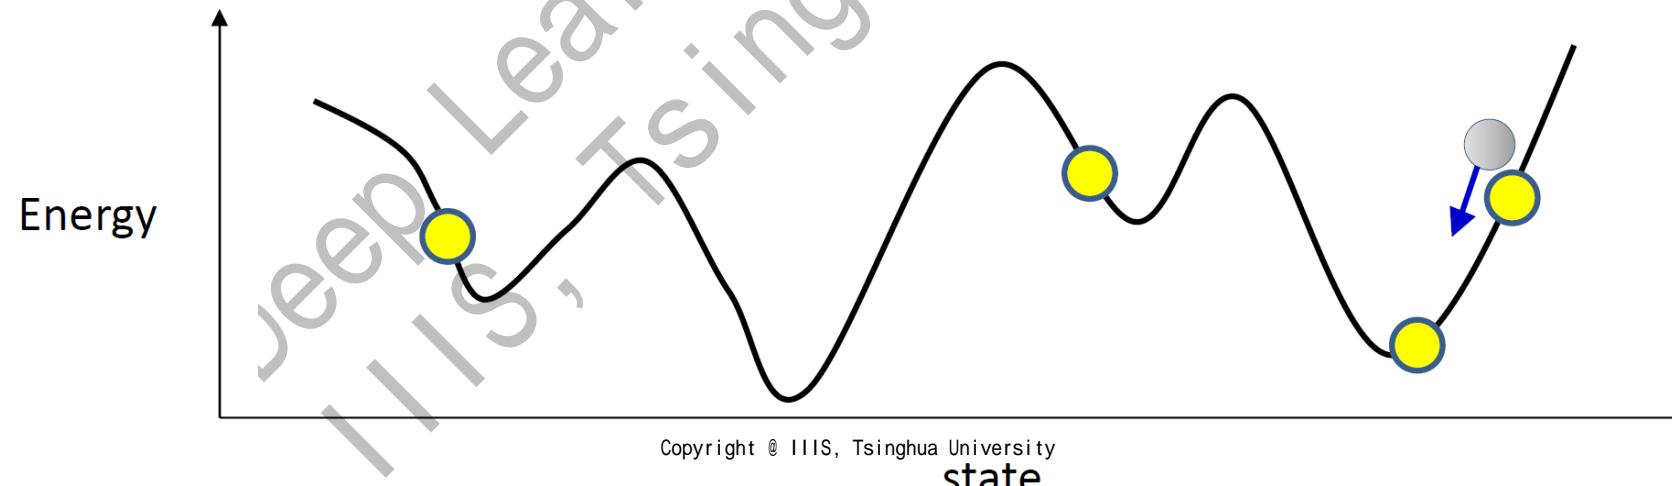
# Hopfield Network: Optimization

- New idea: we only raise the “neighborhood” of desired patterns!
  - It is sufficient to make each desired pattern a valley
  - Note: we want to raise the neighborhood of the decent direction



# Hopfield Network: Optimization

- New idea: we only raise the “neighborhood” of desired patterns!
  - It is sufficient to make each desired pattern a valley
  - Note: we want to raise the neighborhood of the decent direction
- Implementation
  - We initialize  $y'$  by the desired patterns
  - **Only perform evolution for a few steps!**



# Hopfield Network: SGD Optimization

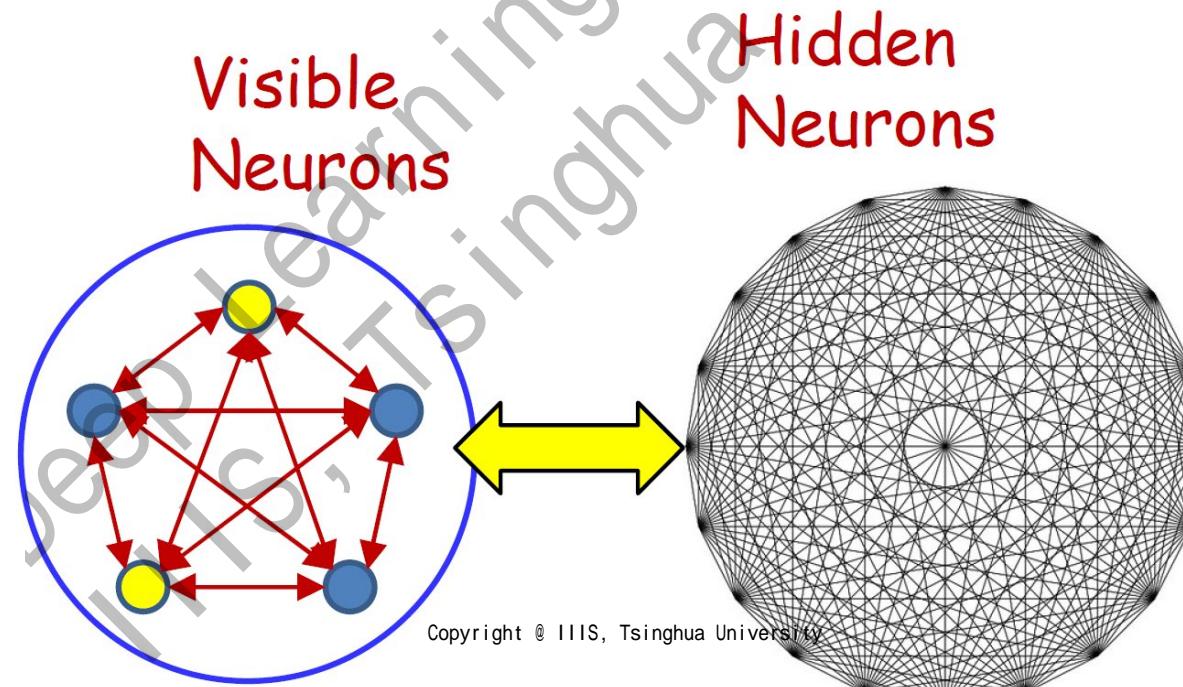
- SGD update rule for  $W$

$$W \leftarrow W - \eta \left( E_{y \in P} [yy^T] - E_{y'} [y'y'^T] \right)$$

- Compute outer-products of random desired pattern  $y$
- Initialize  $y'$  by a random desired pattern**
  - Run evolution for random  $y'$  for a few timesteps (2~4)
  - Calculate outer-product of  $y'$
- Compute gradient and update  $W$
- In theory,  $O(N)$  patterns can be stored in the network (with undesired valleys)
  - How to store more patterns?

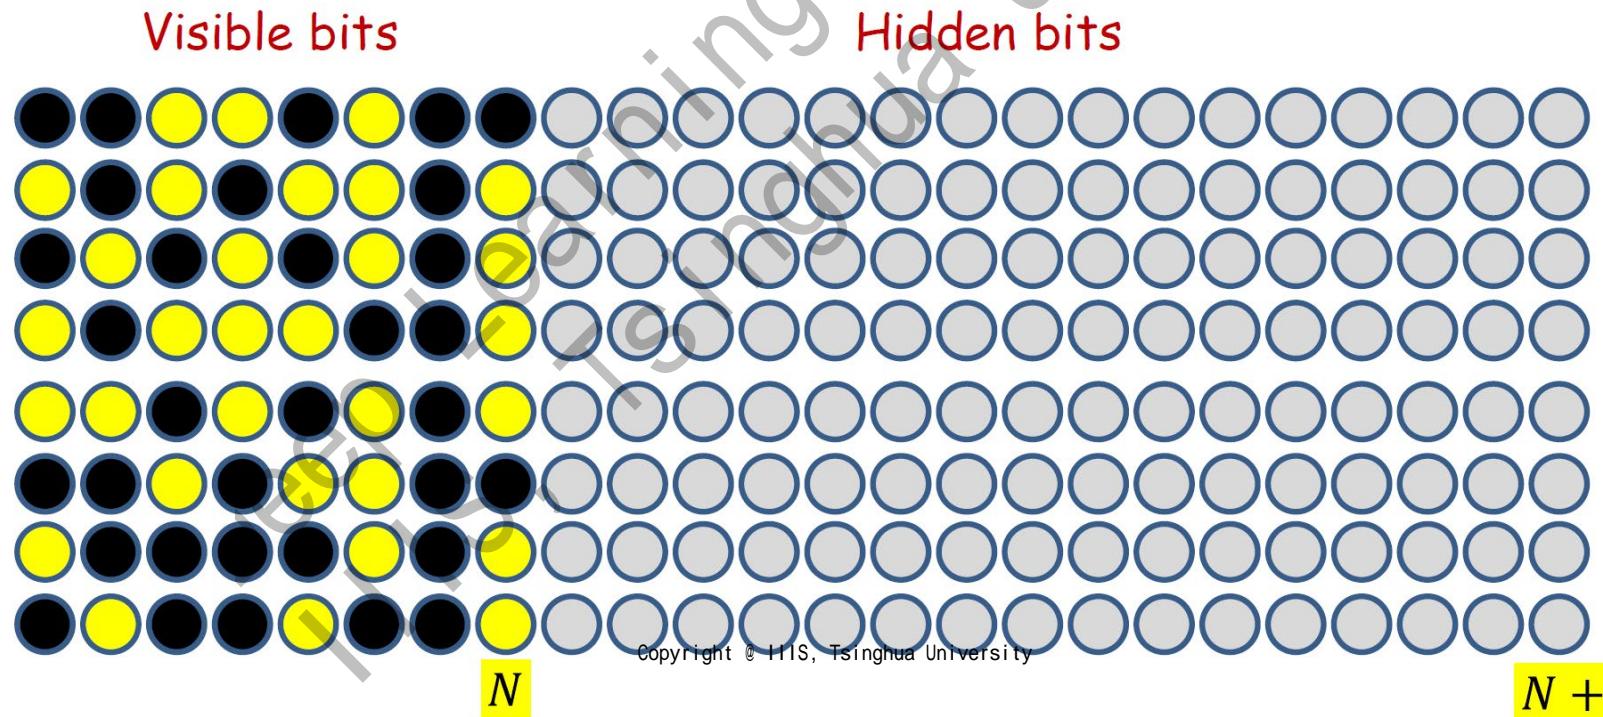
# The Expanded Network

- Idea: introduce redundant neurons to increase network capacity
- Original  $N$  neurons for patterns: visible neurons
- Additional  $K$  neurons: hidden neurons



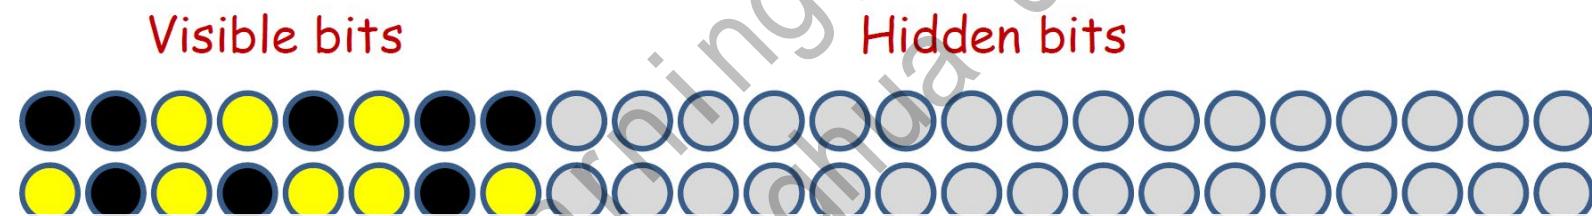
# The Expanded Network

- Idea: introduce redundant neurons to increase network capacity
- Original  $N$  neurons for patterns: visible neurons
- Additional  $K$  neurons: hidden neurons

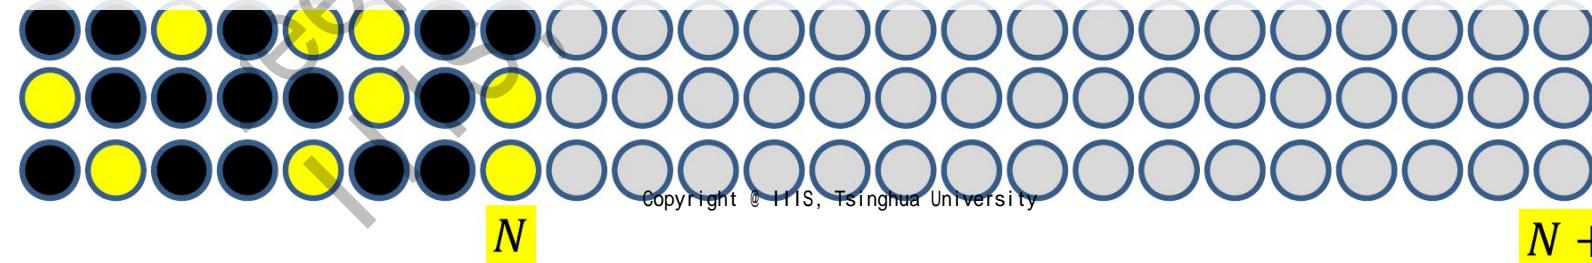


# The Expanded Network

- $N$  dimensional pattern  $\rightarrow N + K$  dimension
  - Q1: How can we store the patterns with  $K$  additional units? (random filling?)
  - Q2: How to retrieve the desired patterns? (perform evolution?)



We will have an elegant solution by converting a Hopfield network to a probabilistic model  $P(v, h)!$

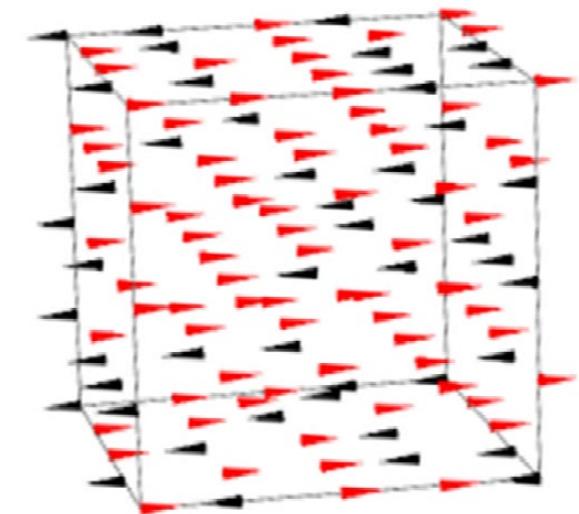


# Today's Lecture: Energy-Based Models

- A particularly flexible and general form of ***generative model***
- Part 1: Hopfield Network
  - The simplest model that can memorize and generate patterns
- Part 2: Boltzmann Machine
  - The first deep generative model
- Part 3: General Energy-Based Models & Sampling Methods

# The Helmholtz Free Energy

- Recap: A thermodynamic (热力学) system
  - A probabilistic system
  - Hopfield network is a simplified deterministic version
- A thermodynamic system at temperature  $T$ 
  - $P_T(S)$  the probability of the system at state  $S$
  - $E_T(S)$  the potential energy at state  $S$
  - $U_T$  the internal energy, the capability to do work
  - $H_T$  the entropy, internal disorder of the system
  - $k$  Boltzmann constant
  - **Free energy**  $F_T = U_T - kT H_T$



# The Helmholtz Free Energy

- Free energy

$$F_T = \sum_S P_T(S) E_T(S) + kT \sum_S P_T(S) \log P_T(S)$$

- Boltzmann distribution (also known as Gibbs distribution)

$$P_T(S) = \frac{1}{Z} \exp\left(-\frac{E_T(S)}{kT}\right)$$

- Minimum Free-Energy Principle: minimize  $F_T$  w.r.t.  $P_T(S)$
- The probability distribution of states at equilibrium
- $Z$  normalizing constant

Given an energy function  $E_T(S)$ , if we follow a proper physical evolution process, the system state will converge to the Boltzmann distribution

# Stochastic Hopfield Network

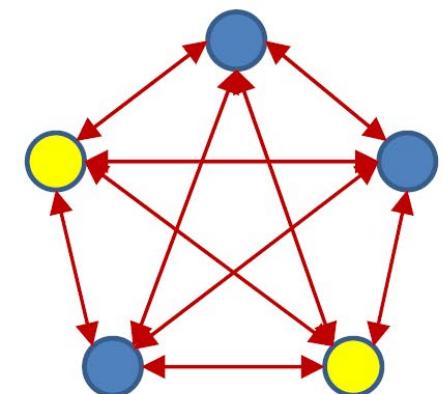
- Let's model our Hopfield network as a thermodynamic system
  - $T = k = 1$  for simplicity
  - Energy

$$E(y) = - \sum_{i < j} w_{ij} y_i y_j - b_i y_i$$

- Boltzmann Probability

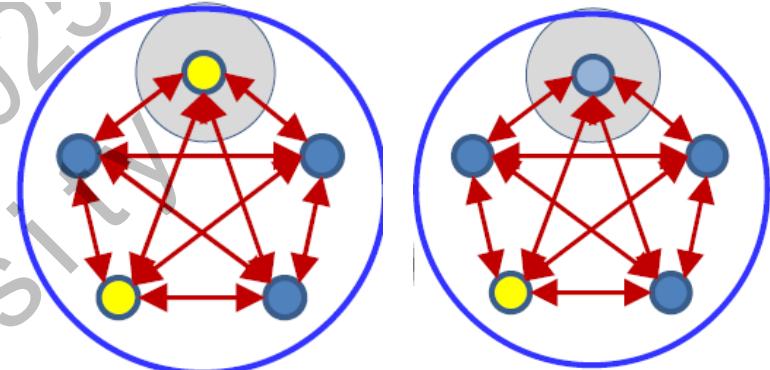
$$P(y) = \frac{1}{Z} \exp(-E(y)) = \frac{1}{Z} \exp\left(\sum_{i < j} w_{ij} y_i y_j + b_i y_i\right)$$

- Stochastic Hopfield Network
  - $P(y)$  models the stationary probability distribution of states  $y$  given  $E(y)$
  - We generate patterns by sampling from  $P(y)$



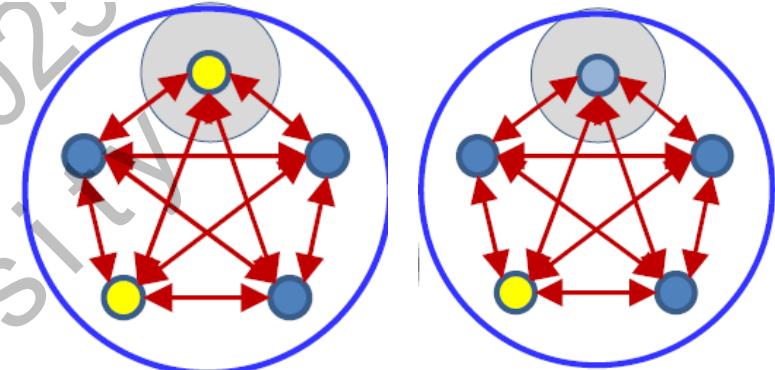
# Stochastic Hopfield Network

- Let's consider the “flip” operation
  - Deterministic → probabilistic
  - Goal: change  $y_i$  to 1 with probability  $P(y_i = 1|y_{j \neq i})$
- Assume  $y$  and  $y'$  only differ at position  $i$  and  $y'_i = -1$ 
  - $\log P(y) = -E(y) + C$
  - $E(y) = -\sum_{i < j} w_{ij} y_i y_j - b_i y_i$
  - $\log P(y) - \log P(y') = E(y') - E(y) = -\sum_j w_{ij} y_j - 2b_i$
  - $\log \frac{P(y)}{P(y')} = \log \frac{P(y_i = 1|y_{j \neq i})P(y_{j \neq i})}{P(y'_i = -1|y'_{j \neq i})P(y'_{j \neq i})} = \log \frac{P(y_i = 1|y_{j \neq i})}{1 - P(y_i = 1|y_{j \neq i})} = -\sum_j w_{ij} y_j - 2b_i$



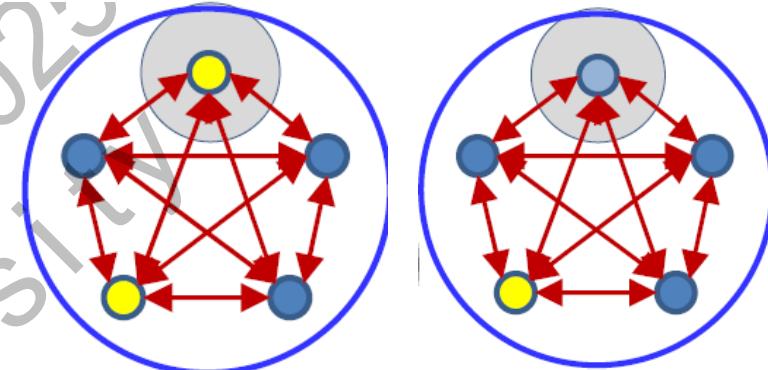
# Stochastic Hopfield Network

- Let's consider the “flip” operation
  - Deterministic → probabilistic
  - Goal: change  $y_i$  to 1 with probability  $P(y_i = 1 | y_{j \neq i})$
- Assume  $y$  and  $y'$  only differ at position  $i$  and  $y'_i = -1$ 
  - $\log P(y) = -E(y) + C$
  - $E(y) = -\sum_{i < j} w_{ij}y_i y_j - b_i y_i$
  - $\log P(y) - \log P(y') = E(y') - E(y) = -\sum_j w_{ij}y_j - 2b_i$
  - $\log \frac{P(y)}{P(y')} = \log \frac{P(y_i = 1 | y_{j \neq i})P(y_{j \neq i})}{P(y'_i = -1 | y'_{j \neq i})P(y'_{j \neq i})} = \log \frac{P(y_i = 1 | y_{j \neq i})}{1 - P(y_i = 1 | y_{j \neq i})} = -\sum_j w_{ij}y_j - 2b_i$
  - A sigmoid conditional:  $P(y_i = 1 | y_{j \neq i}) = \frac{1}{1 + \exp(-\sum_j w_{ij}y_j - 2b_i)}$



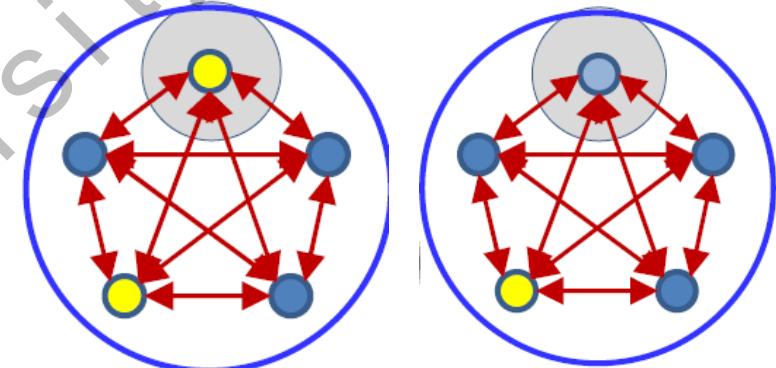
# Stochastic Hopfield Network

- The whole update rule
  - Field at  $y_i$ :  $z_i = \sum_j w_{ij}y_j + 2b_i$
  - $P(y_i = 1 | y_{j \neq i}) = \frac{1}{1 + \exp(-z_i)} = \sigma(z_i)$
- Evolving the network
  - Randomly initialize  $y$
  - Cycle over  $y_i$ , fixed other variables fixed and sample  $y_i$  according to the conditional probability
  - After “convergence”, we can get samples of  $y$  according to  $P(y)$
  - *This sampling procedure is called Gibbs sampling*
  - **How can we retrieve a stored pattern???**
    - **This is a stochastic process!**



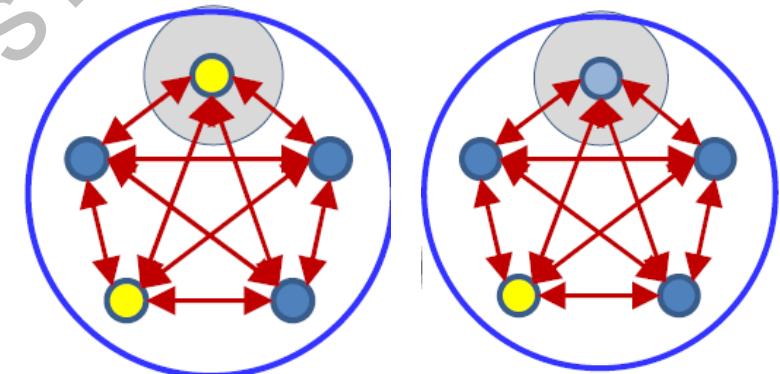
# Stochastic Hopfield Network

- Network evolution
  - initialize  $y_0$
  - For  $1 \leq i \leq N$ ,  $y_i(t + 1) \sim \text{Bernoulli}(\sigma(z_i(t)))$
  - Until convergence
- Retrieve a stored (low energy / high probability) pattern  $y$ 
  - Given sequence of samples  $y_0, \dots, y_L$
  - Simply take the average of final  $M$  samples
$$y_i = I \left[ \frac{1}{M} \sum_{t=L-M+1}^L y_i(t) > 0 \right]$$
  - If you want a probability instead of a single vector, you can use the frequency derived from  $\{y_{L-M+1}, \dots, y_L\}$  to approximate the stationary distribution
  - **In many applications, we simply take  $M = 1$  (output  $y_L$ )**



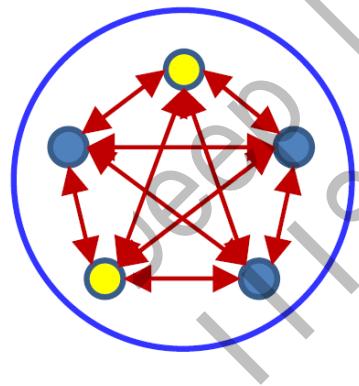
# Stochastic Hopfield Network: Annealing

- Find the state with lowest energy?
- Network evolution with temperature annealing
  - initialize  $y_0, T \leftarrow T_{\max}$
  - Repeat
    - Repeat a few cycles
      - For  $1 \leq i \leq N, y_i(T) \sim \text{Bernoulli} \left( \sigma \left( \frac{1}{T} z_i(T) \right) \right)$
      - $y_i(\alpha T) \leftarrow y_i(T); T \leftarrow \alpha T$
    - Until convergence
  - Final state as the retrieved pattern
    - With temperature annealing, the system will converge to the most likely state
    - Possibly local minimum in practice



# Boltzmann Machine

- A generative Model (simplified)
  - $E(y) = -\frac{1}{2}y^T W y$
  - $P(y) = \frac{1}{Z} \exp\left(-\frac{E(y)}{T}\right)$
  - Parameter  $W$
- It has a probability for producing any binary pattern  $y$ 
  - We assume  $y_i = 0$  or  $1$  (or  $\pm 1$ )



How to learn  $W$  for desired patterns?

$$\mathbf{z}_i = \frac{1}{T} \sum_j w_{j,i} y_j$$

$$P(y_i = 1 | y_{j \neq i}) = \frac{1}{1 + e^{-z_i}}$$

# Boltzmann Machine: Training

- Goal
  - Remember a set of desired patterns  $P = \{y^p\}$
  - Now we have a probability distribution  $P(y)$  with parameter  $W$
- Objective: maximum likelihood learning (assume  $T = 1$ )
  - Probability of a particular pattern

$$P(y) = \frac{\exp\left(\frac{1}{2}y^T W y\right)}{\sum_{y'} \exp\left(\frac{1}{2}y'^T W y'\right)}$$

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp\left(\frac{1}{2}y'^T W y'\right)$$

# Boltzmann Machine: Training

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp\left(\frac{1}{2} y'^T W y'\right)$$

- Gradient Ascent  $\nabla_{w_{ij}} L$

# Boltzmann Machine: Training

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp\left(\frac{1}{2} y'^T W y'\right)$$

- Gradient Ascent  $\nabla_{w_{ij}} L$

- $\bullet \nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j$

# Boltzmann Machine: Training

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp\left(\frac{1}{2} y'^T W y'\right)$$

- Gradient Ascent  $\nabla_{W_{ij}} L$

$$\nabla_{W_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \sum_{y'} \frac{\exp\left(\frac{1}{2} y'^T W y'\right)}{Z} \cdot y'_i y'_j$$

Exponentially many terms!

# Boltzmann Machine: Training

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp\left(\frac{1}{2} y'^T W y'\right)$$

- Gradient Ascent  $\nabla_{w_{ij}} L$

- $\nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \sum_{y'} \frac{\exp\left(\frac{1}{2} y'^T W y'\right)}{Z} \cdot y'_i y'_j$

- $\nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \mathbb{E}_{y'} [y'_i y'_j]$  **Monte-Carlo Approximation**

- Draw a set of samples  $S$  for  $y'$  according to the probability,

- $\nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \frac{1}{|S|} \sum_{y' \in S} y'_i y'_j$

# Boltzmann Machine: Training

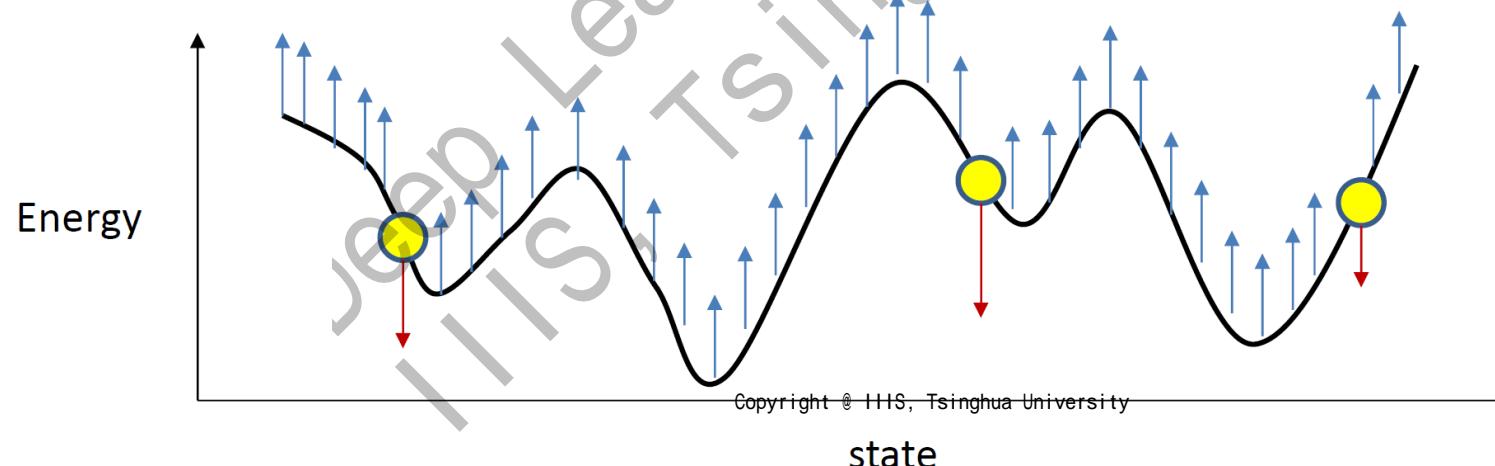
- Maximize log-likelihood with  $M$  Monte-Carlo samples

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \frac{1}{M} \sum_{y' \in S} y'_i y'_j$$

- How to draw samples from  $P(y)$ ?
  - Running the stochastic network (Gibbs sampling)
    - Randomly initialize  $y(0)$
    - Cycle over  $y_i(t)$ , sampling according to  $P(y_i(t)|y_{j \neq i}(t))$
    - After convergence, we get a sequence of samples  $\{y(0), \dots, y(L)\}$
    - Get the final  $M$  states as samples  $S = \{y(L - M + 1), \dots, y(L)\}$

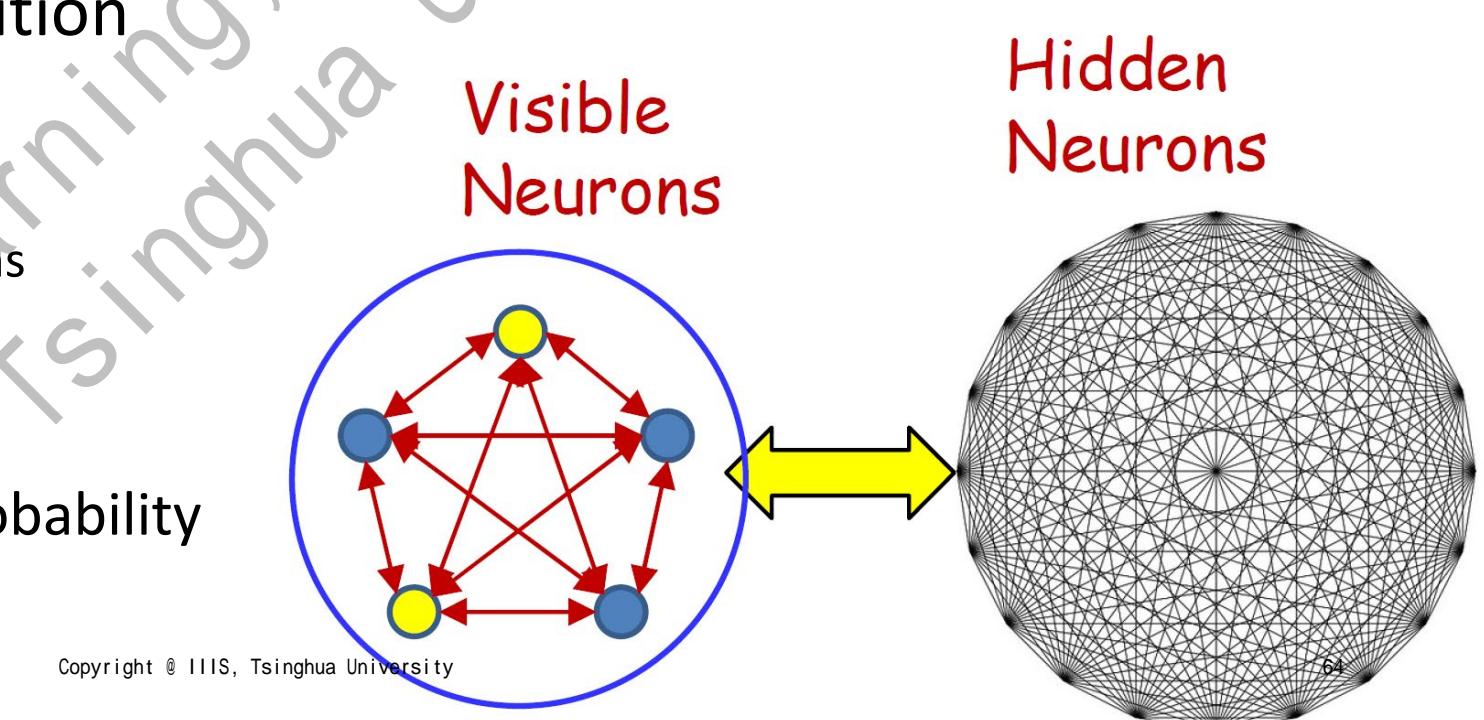
# Boltzmann Machine: Training

- Overall Training
  - Initialize  $W$
  - Maximize log-likelihood with  $M$  Monte-Carlo samples
- $w_{ij} \leftarrow w_{ij} + \eta \nabla_{w_{ij}} L(W)$  (*we are maximizing likelihood*)



# Boltzmann Machine with Hidden Neurons

- Let's get back to hidden neurons!
  - $v$  visible neurons (visible patterns),  $h$  hidden neurons (latent variables)
  - $y = (v, h)$
- A joint probability distribution
  - $P(y) = P(v, h)$
  - $P(v) = \sum_h P(v, h)$ 
    - We only care about patterns
    - **The marginal distribution!**
- New objective
  - Maximize the marginal probability



# Boltzmann Machine with Hidden Neurons

- Maximum log-likelihood learning

$$P(v) = \sum_h P(v, h) = \sum_h \frac{\exp(y^T W y)}{\sum_{y'} \exp(y'^T W y')}$$
$$L(W) = \frac{1}{|P|} \sum_{v \in P} \log \left( \sum_h \exp(y^T W y) \right) - \log \left( \sum_{y'} \exp(y'^T W y') \right)$$

- Gradient  $\nabla L(W)$ ?

# Boltzmann Machine with Hidden Neurons

- Maximum log-likelihood learning

$$P(v) = \sum_h P(v, h) = \sum_h \frac{\exp(y^T W y)}{\sum_{y'} \exp(y'^T W y')}$$
$$L(W) = \frac{1}{|P|} \sum_{v \in P} \log \left( \sum_h \exp(y^T W y) \right) - \log \left( \sum_{y'} \exp(y'^T W y') \right)$$

- Gradient  $\nabla L(W)$ ?

**Monte-Carlo Estimate!**

# Boltzmann Machine with Hidden Neurons

- Maximum log-likelihood learning

$$P(v) = \sum_h P(v, h) = \sum_h \frac{\exp(y^T W y)}{\sum_{y'} \exp(y'^T W y')}$$
$$L(W) = \frac{1}{|P|} \sum_{v \in P} \log \left( \sum_h \exp(y^T W y) \right) - \log \left( \sum_{y'} \exp(y'^T W y') \right)$$

- Gradient  $\nabla L(W)$ ?
  - The first term is also in the form of log-sum
  - Monte Carlo estimates for each  $v \in P$ !

# Boltzmann Machine with Hidden Neurons

- Maximum log-likelihood learning

$$\nabla_{w_{ij}} L(W) = \frac{1}{|P|} \sum_{v \in P} E_h[y_i y_j] - E_{y'}[y'_i y'_j]$$

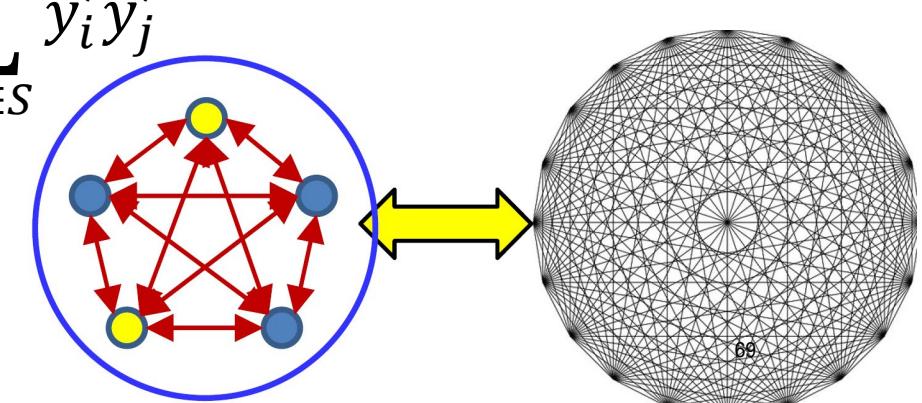
- Second term
  - Freely generate samples w.r.t.  $p(y)$
  - Random initialization, cyclic Gibbs sampling
- First term
  - Generate samples w.r.t.  $p(y)$  conditioned on a fixed  $v$
  - Randomly initialize  $h$ , run Gibbs sampling over  $h$

# Boltzmann Machine with Hidden Neurons

- Overall Training
  - Initialize  $W$
  - For  $v \in P$ , fix the visible neurons, run Gibbs sampling to get  $K$  samples
    - Collect all conditioned samples as  $S_c$
  - Randomly initialize all neurons, run Gibbs sampling to get  $M$  samples
    - Collect free samples as  $S$
  - Maximize log-likelihood with  $N_p K + M$  Monte-Carlo samples

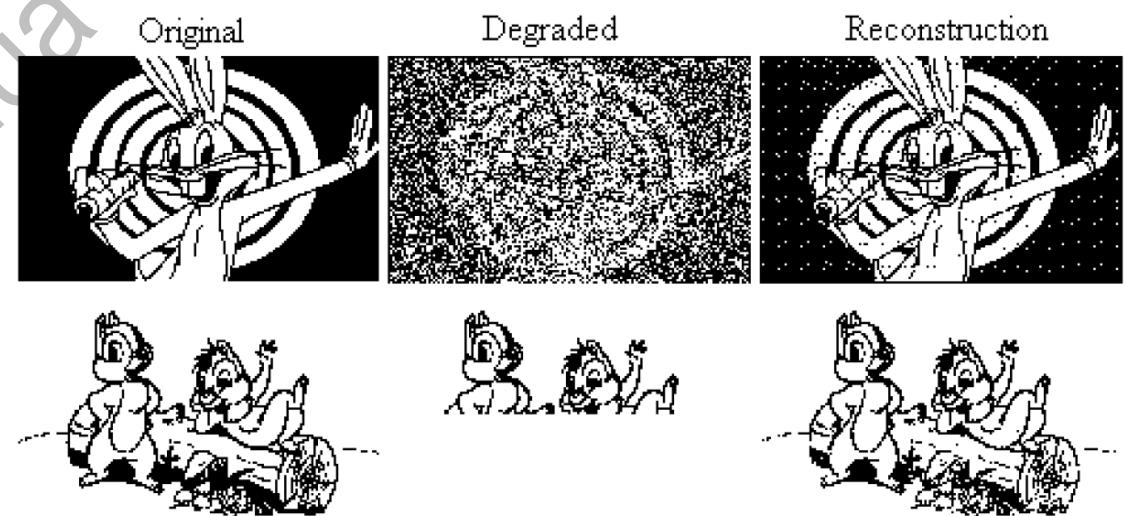
$$\nabla_{w_{ij}} L(W) = \frac{1}{N_p K} \sum_{y \in S_c} y_i y_j - \frac{1}{M} \sum_{y' \in S} y'_i y'_j$$

$$\bullet w_{ij} \leftarrow w_{ij} + \eta \nabla_{w_{ij}} L(W)$$



# Boltzmann Machine

- Summary
  - A stochastic version of Hopfield Network
    - Nice mathematical properties
    - Large capacity for storing patterns (with hidden neurons)
  - Pattern generation
    - Gibbs sampling
  - Pattern completion
    - Conditioned Gibbs sampling
  - **Classification??**
    - $y = (v, h, c)$ ,  $c$  is label
    - $c$  as a one-hot vector (0-1 variables)
    - Posterior  $P(c|v)$
    - Even conditional generation:  $P_{\text{Copy}}(v|c)$ !



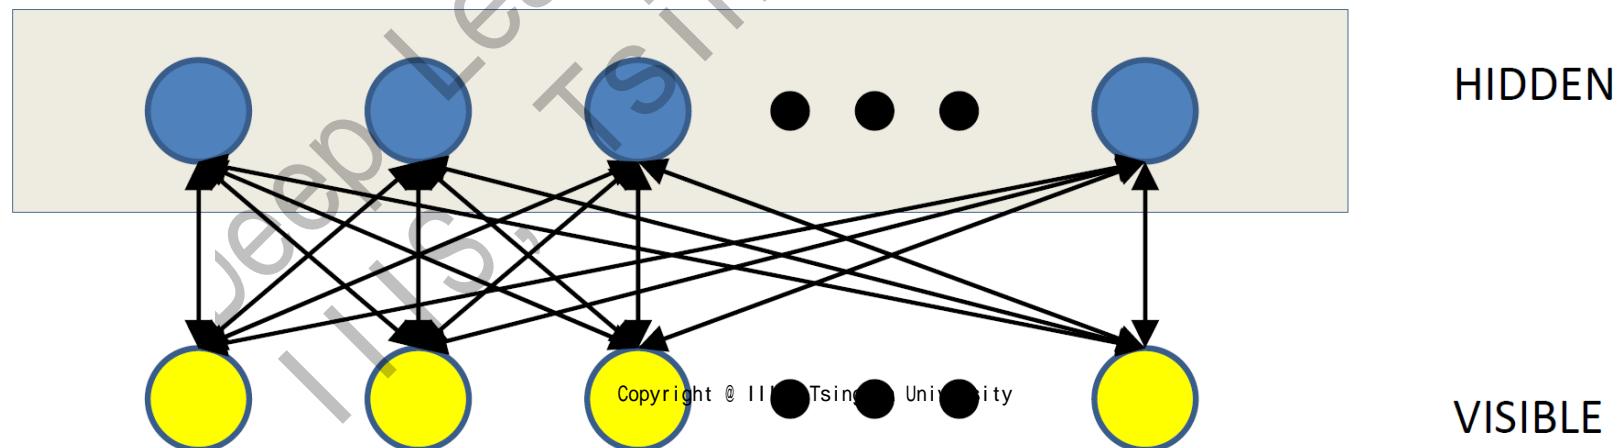
Hopfield network reconstructing degraded images  
from noisy (top) or partial (bottom) cues.

# Boltzmann Machine

- The issue
  - Training is hard!
  - Gibbs sampling may take a very long time to converge
    - also called *mixing-time*
  - Not really applicable for large problems
- Can we design a better structure for faster Gibbs sampling mixing?

# Restricted Boltzmann Machine

- A particularly structured Boltzmann Machine
  - A partitioned structure
  - Hidden neurons are only connected to visible neurons
  - No intra-layer connections
  - *Invented under the name Harmonium by Paul Smolensky in 1986*
  - *Became promise after Hinton invented fast learning algorithms in mid-2000*



# Restricted Boltzmann Machine

- Computation Rules: same as Boltzmann machine

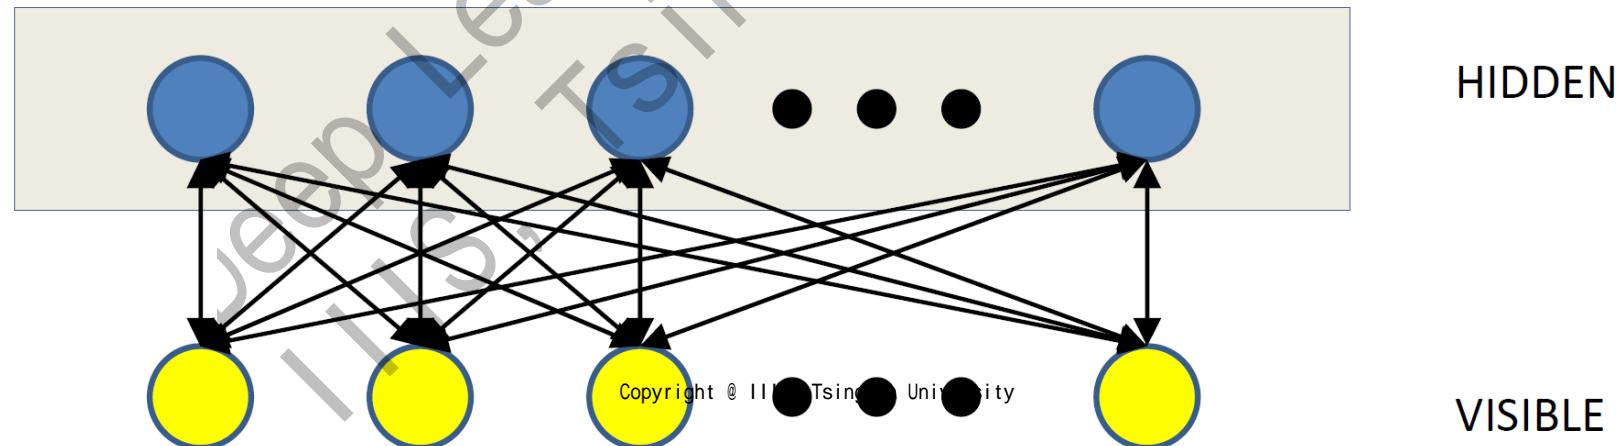
- Hidden neurons  $h_i$

$$z_i = \sum_j w_{ij} v_j, \quad P(h_i = 1 | v_j) = \frac{1}{1 + \exp(-z_i)}$$

- Visible neurons  $v_j$

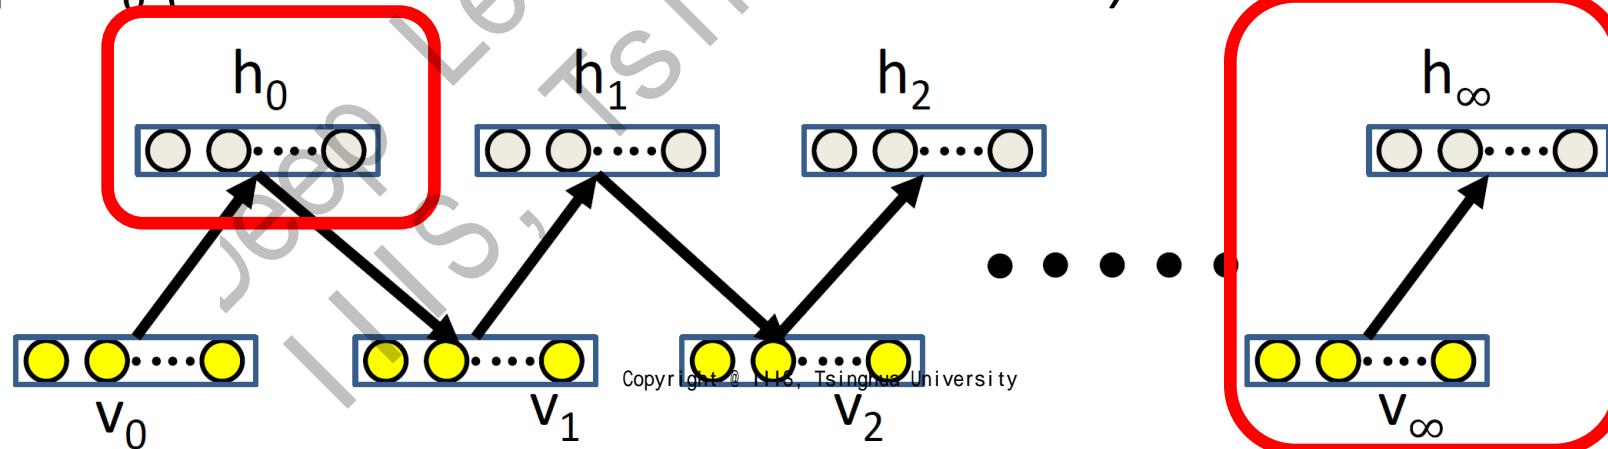
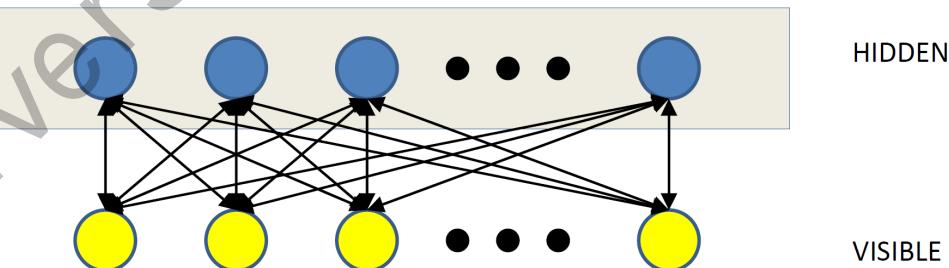
$$z_j = \sum_i w_{ij} h_i, \quad P(v_j = 1 | h_i) = \frac{1}{1 + \exp(-z_j)}$$

↔ Iterative Sampling!



# Restricted Boltzmann Machine

- Sampling
  - Randomly initialize visible neurons  $v_0$
  - Iterative between hidden and visible neurons
  - Get final sample  $(v_\infty, h_\infty)$
- Conditioned sampling?
  - Initialize  $v_0$  as the desired pattern
  - Sample  $h_0$  (*the conditional distribution is exact!*)

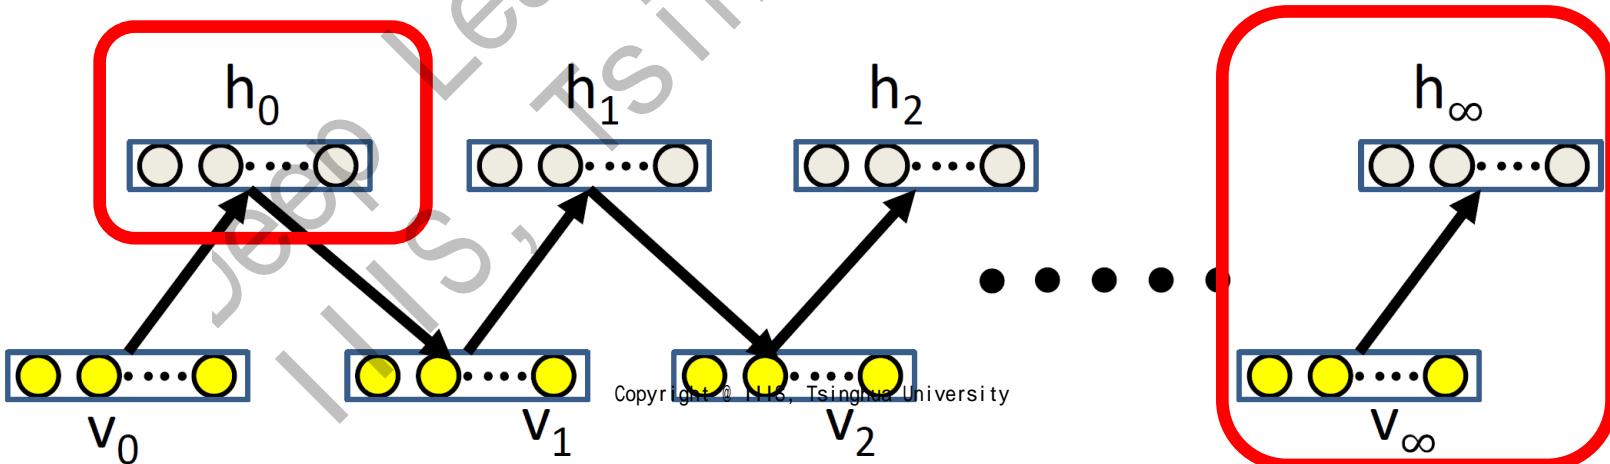


# Restricted Boltzmann Machine

- Maximum Likelihood Estimate

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P K} \sum_{v \in P} v_{0i} h_{0j} - \frac{1}{M} \sum_{v \in M} v_{\infty i} h_{\infty j}$$

- No need to lift up the entire energy landscape ... (recap)

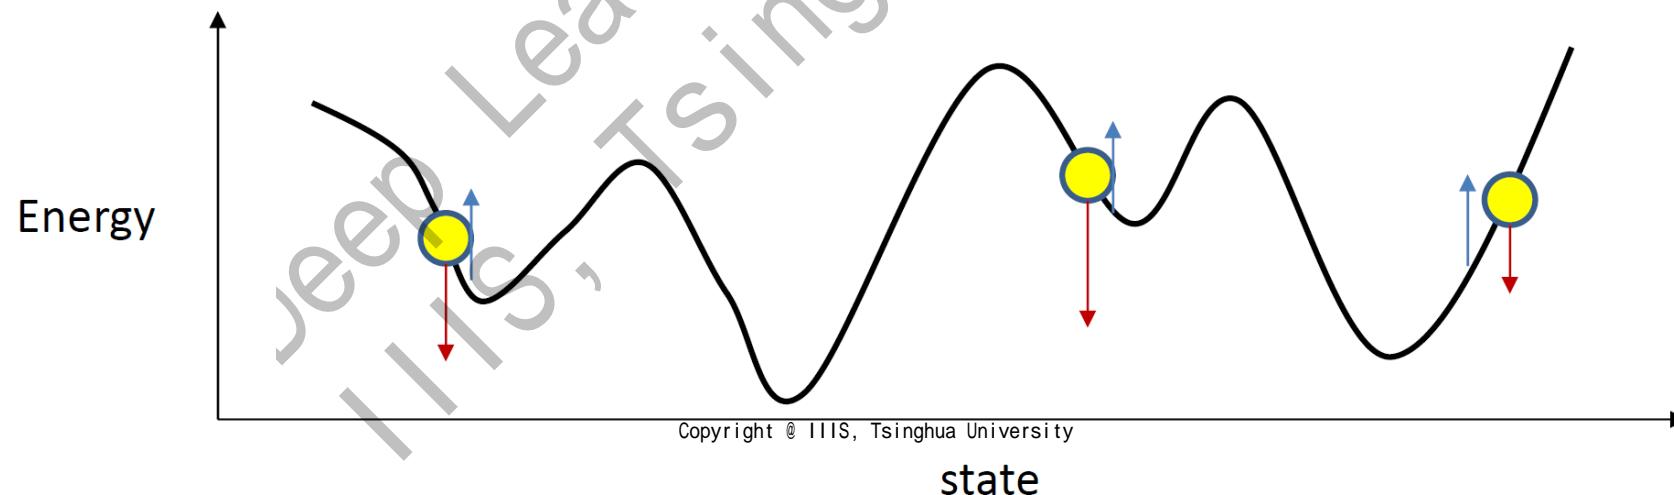


# Restricted Boltzmann Machine

- Maximum Likelihood Estimate

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P K} \sum_{v \in P} v_{0i} h_{0j} - \frac{1}{M} \sum v_{\infty i} h_{\infty j}$$

- We can start sampling with a given  $v_0$ 
  - Raising the neighborhood of the desired patterns will be sufficient

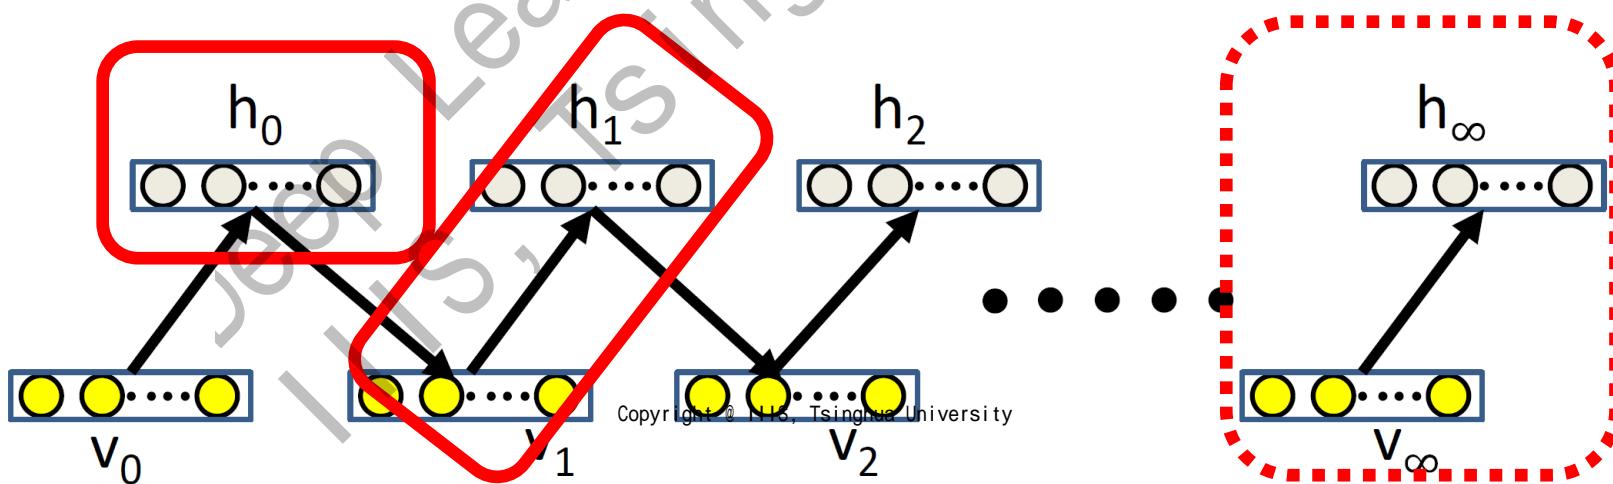


# Restricted Boltzmann Machine

- Maximum Likelihood Estimate

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P K} \sum_{v \in P} v_{0i} h_{0j} - \frac{1}{M} \sum v_{\infty i} h_{\infty j}$$

- Directly run Gibbs sampling from  $v_0$  for 3 steps will be sufficient!



# Restricted Boltzmann Machine

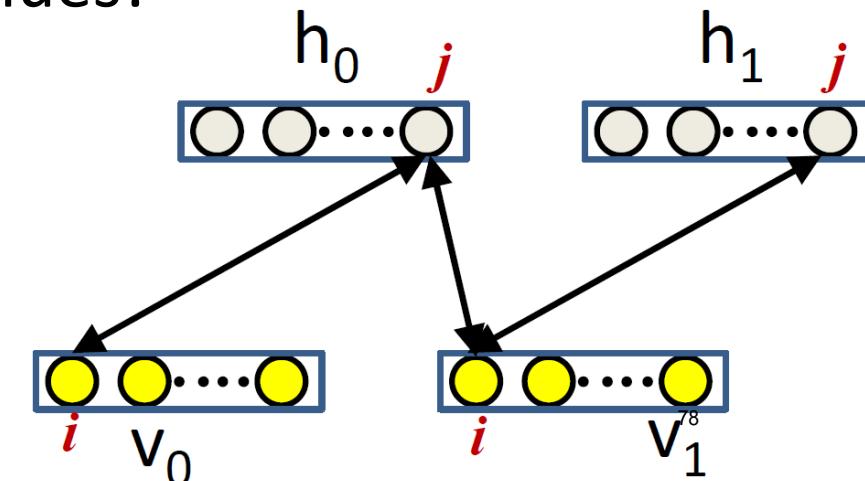
- Maximum Likelihood Estimate

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P} \sum_{v \in P} v_{0i} h_{0j} - v_{1i} h_{1j}$$

- Only 3 Gibbs sampling steps are needed!

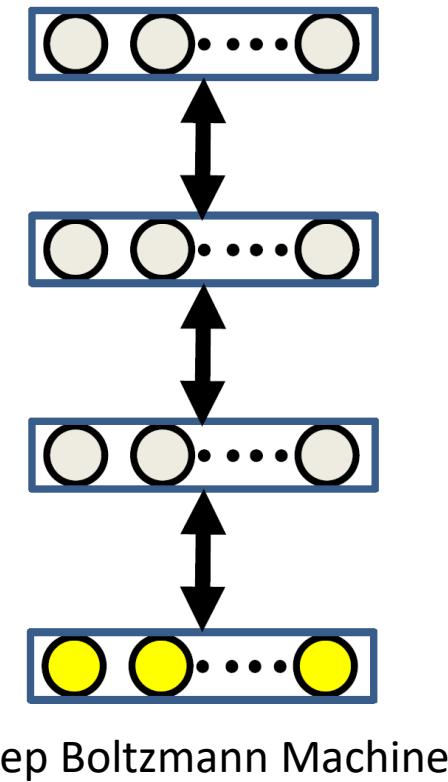
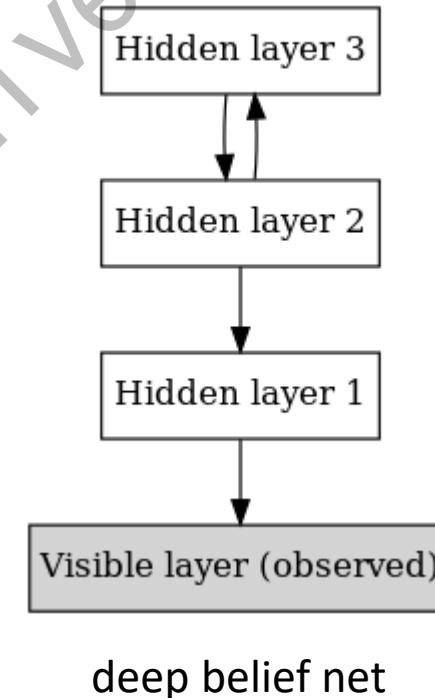
- We can also extend (R)BMs to continuous values!

- If we can explicitly sample from  $P(y_i | y_{j \neq i})$
- Exponential family! (FYI 😊)
  - “Exponential Family Harmoniums with an Application to Information Retrieval”, Welling et al., 2004



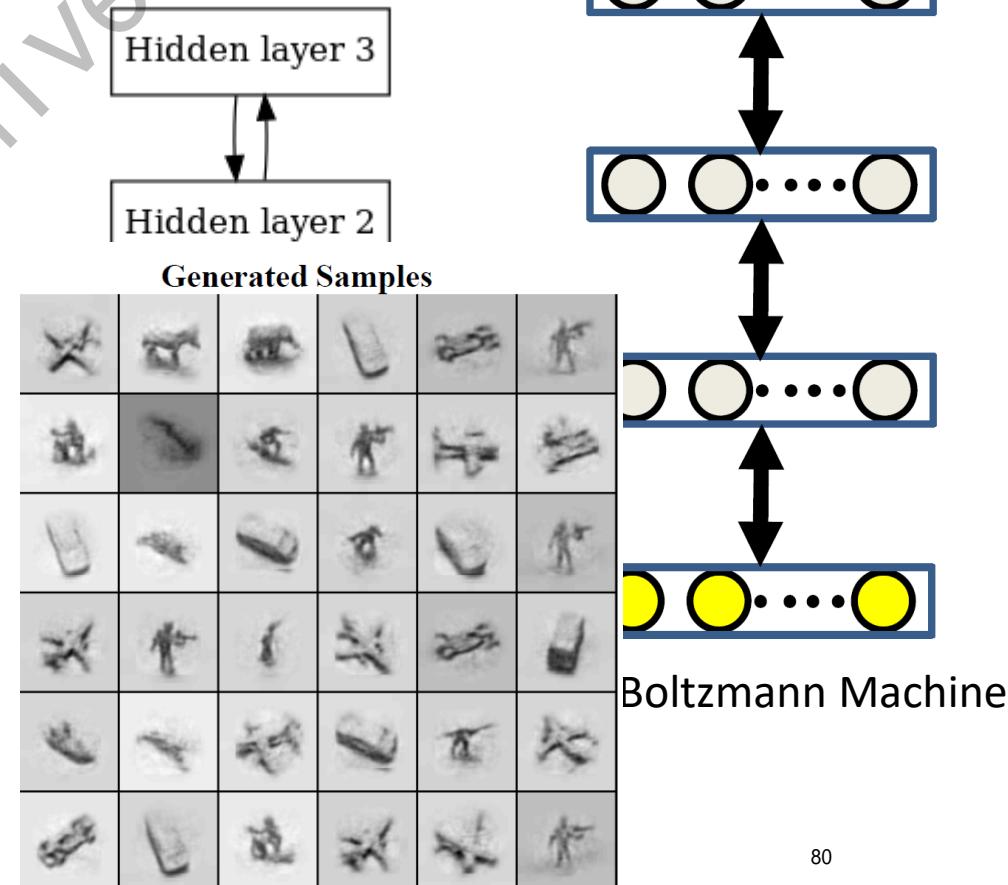
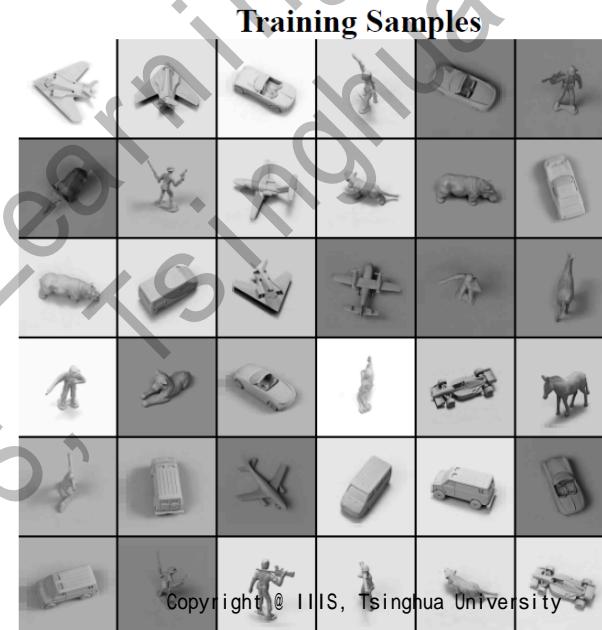
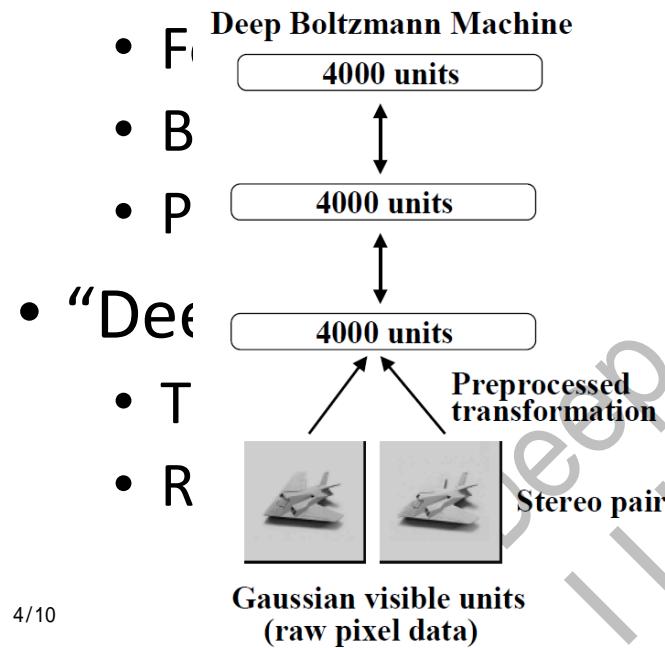
# Deep Boltzmann Machine

- Can we have a ***deep*** version of RBM?
  - Deep Belief Net (2006)
  - Deep Boltzmann Machine (2009)
- Sampling?
  - Forward pass: bottom-up
  - Backward pass: top-down
  - Practical Trick: Layer-by-layer pretraining
- “Deep Boltzmann Machine”, AISTATS 2009
  - The very first deep generative model
  - Ruslan Salakhutdinov & Geoffrey Hinton



# Deep Boltzmann Machine

- Can we have a ***deep*** version of RBM?
  - Deep Belief Net (2006)
  - Deep Boltzmann Machine (2009)
- Sampling?



# Nobel Prize in Physics 2024



© Nobel Prize Outreach. Photo:  
Nanaka Adachi

**John J. Hopfield**

Prize share: 1/2



© Nobel Prize Outreach. Photo:  
Clément Morin

**Geoffrey Hinton**

Copyright © IIIS, Tsinghua University

Prize share: 1/2

# Today's Lecture: Energy-Based Models

- A particularly flexible and general form of *generative model*
- Part 1: Hopfield Network
  - The simplest model that can memorize and generate patterns
- Part 2: Boltzmann Machine
  - The first deep generative model
- Part 3: General Energy-Based Models & Sampling Methods

# Energy-Based Model

- Goal of generative model
  - A probability distribution of “patterns”  $P(x)$
- Requirement
  - $P(x) \geq 0$  (non-negative)
  - $\int_x P(x)dx = 1$  (sum to 1)
- Energy-Based Model
  - Energy function:  $E(x; \theta)$  parameterized by  $\theta$
  - $P(x) = \frac{1}{Z} \exp(-E(x; \theta))$
  - $Z = \int_x \exp(-E(x; \theta)) dx$  *partition function*

Why use  $\exp()$  function?  
e.g.  $|x|$  or  $|x|^2$

# Energy-Based Model

- A particular class of density function

$$P(x) = \frac{1}{Z} \exp(-E(x; \theta))$$

- Pros

- Common in statistical physics
- Compatible with log-probability measure to capture large variations
- Exponential family (e.g., Gaussian)
- Extremely flexible, i.e., use any  $E(x)$  you like (e.g., any  $f(x): \mathbb{R}^d \rightarrow \mathbb{R}$ , even CNNs)

- Cons

- Non-trivial to sample and train due to the partition function  $Z$

# Energy-Based Model: Training

- A particular class of density function

$$P(x) = \frac{1}{Z} \exp(-E(x; \theta))$$

- Maximum Likelihood Training

- $L(\theta) = \log P(x) = -E(x; \theta) - \log Z(\theta)$
  - Monte-Carlo estimates for partition function  $Z(\theta)$

- Contrastive Divergence Algorithm

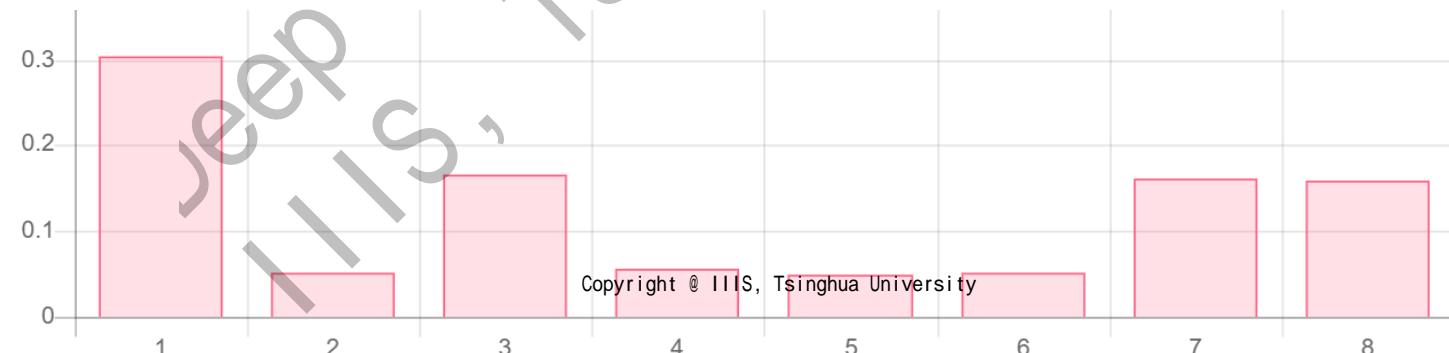
- $\nabla_{\theta} L(\theta) \approx \nabla_{\theta} (-E(x_{train}; \theta) + E(x_{sample}; \theta))$
  - Generating samples is the foundation for both training and generation!

- How to sample from an general energy-based model?

- Or in general: sample from an arbitrary distribution  $p(x)$

# Sampling Methods

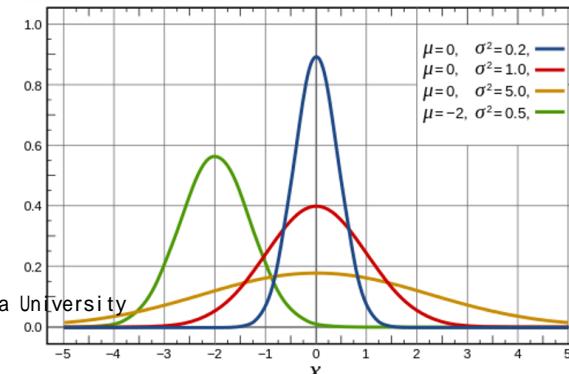
- Goal: sampling from  $P(x)$ 
  - Assume we have a valid probability measure
  - $P(x)$  can be arbitrarily complex (e.g., high-dimensional, continuous, etc)
- Let's start from an easy example
  - Categorical distribution?
  - Solution: uniform sampling, find the category with cumulative density
    - *The mapping from CDF to value is called Inverse distribution function (quantile function)*



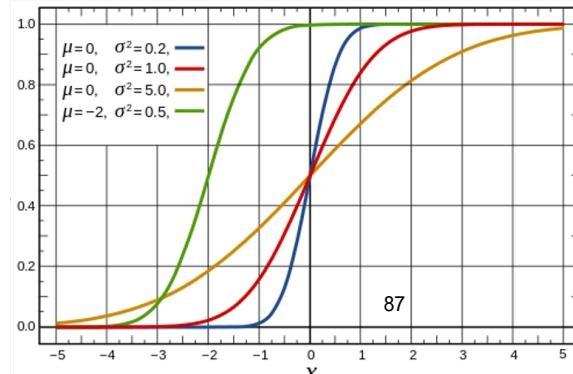
# Sampling Methods

- Goal: sampling from  $P(x)$ 
  - Assume we have a valid probability measure
  - $P(x)$  can be arbitrarily complex (e.g., high-dimensional, continuous, etc)
- Let's start from an easy example
  - Categorical distribution
  - Gaussian distribution?
    - No closed-form CDF!
    - Central-limit theorem
      - Sample  $X_i \sim Beroulli(0.5)$
      - $E[X_i] = 0.5; Var[X_i] = 0.5^2$
      - $S_N = \frac{1}{N} \sum_{i=1}^N X_i$
      - As  $N \rightarrow \infty, \sqrt{N}(S_N - 0.5) \sim N(0, 0.5^2)$

Probability Density Function



Cumulative Density Function



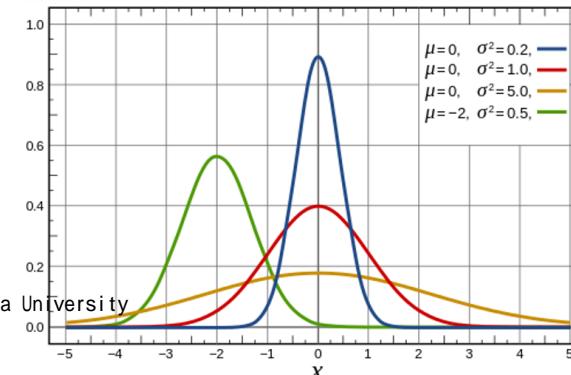
# Sampling Methods

- Goal: sampling from  $P(x)$ 
  - Assume we have a valid probability measure
  - $P(x)$  can be arbitrarily complex (e.g., high-dimensional, continuous, etc)
- Let's start from an easy example
  - Categorical distribution
  - Gaussian distribution?
    - No closed-form CDF!
    - Central-limit theorem
    - **Box–Muller transform**
      - Most practical method (FYI)
      - Uniform  $\rightarrow$  Normal
      - Polar form transformation

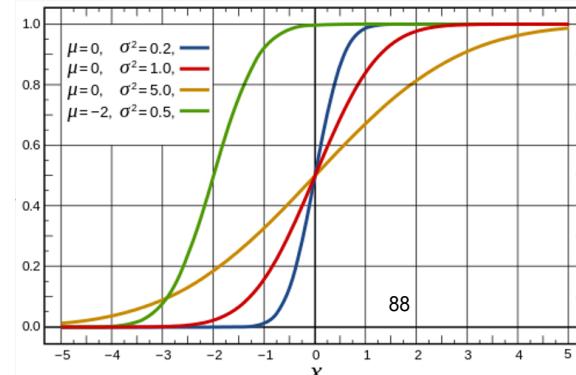
```
def box_muller():
    # Avoid getting u == 0.0
    u1, u2 = 0.0, 0.0
    while u1 < epsilon or u2 < epsilon:
        u1 = random.random()
        u2 = random.random()

    n1 = math.sqrt(-2 * math.log(u1)) * math.cos(2 * math.pi * u2)
    n2 = math.sqrt(-2 * math.log(u1)) * math.sin(2 * math.pi * u2)
    return n1, n2
```

Probability Density Function



Cumulative Density Function



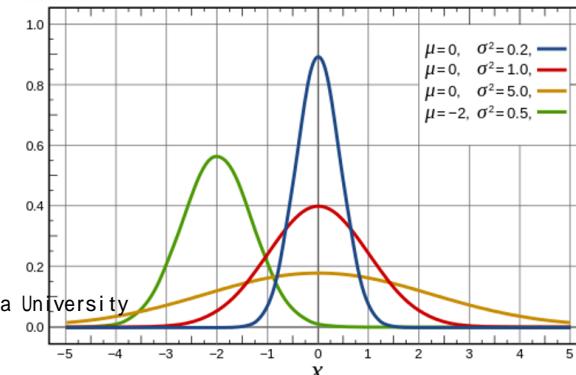
# Sampling Methods

- Goal: sampling from  $P(x)$ 
  - Assume we have a valid probability measure
  - $P(x)$  can be arbitrarily complex (e.g., high-dimensional, continuous, etc)
- Let's start from an easy example
  - Categorical distribution
  - Gaussian distribution?
    - No closed-form CDF!
    - Central-limit theorem
    - Box–Muller transform
    - General case  $x \sim N(\mu, \sigma^2)$
    - High-dimensional case  $x \sim N(\mu, \Sigma)$ 
      - $z \sim N(0, I)$
      - $x = \Sigma z + \mu$

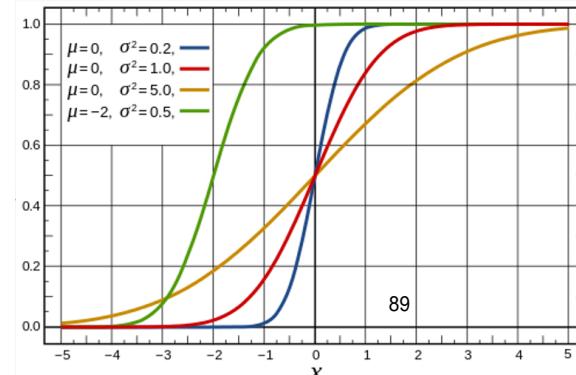
```
def box_muller():
    # Avoid getting u == 0.0
    u1, u2 = 0.0, 0.0
    while u1 < epsilon or u2 < epsilon:
        u1 = random.random()
        u2 = random.random()

    n1 = math.sqrt(-2 * math.log(u1)) * math.cos(2 * math.pi * u2)
    n2 = math.sqrt(-2 * math.log(u1)) * math.sin(2 * math.pi * u2)
    return n1, n2
```

Probability Density Function

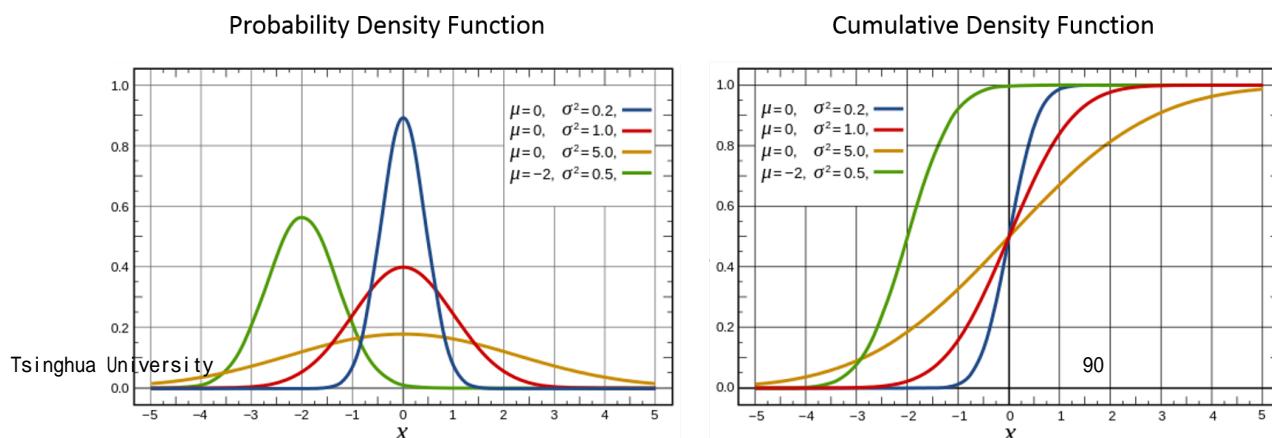
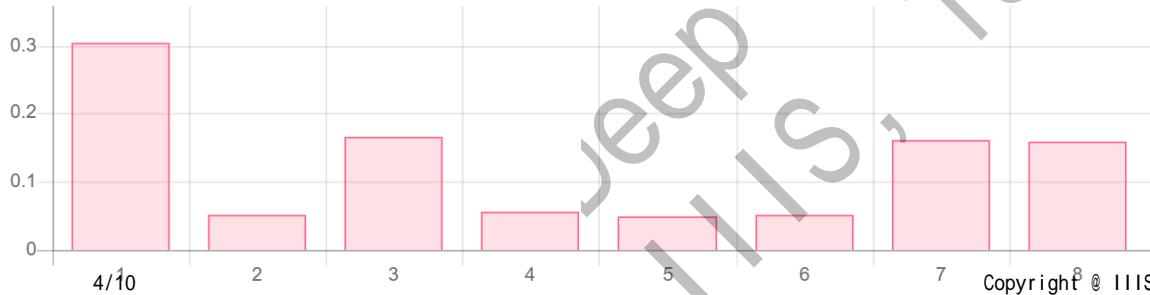


Cumulative Density Function



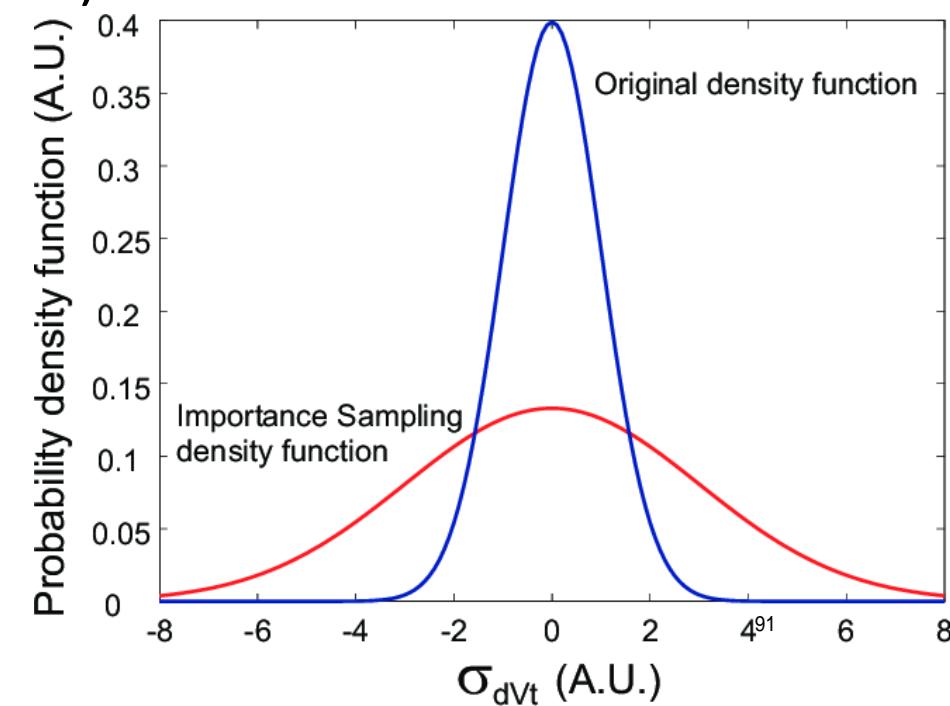
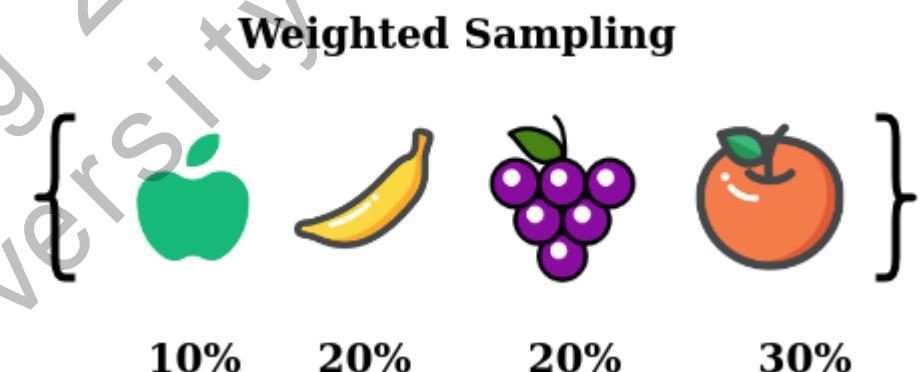
# Sampling Methods

- Goal: sampling from  $P(x)$ 
  - Assume we have a valid probability measure
  - $P(x)$  can be arbitrarily complex (e.g., high-dimensional, continuous, etc)
- Let's start from an easy example
  - Categorical distribution
  - Gaussian distribution
    - Idea: (1) use “easy” distributions to draw sample & (2) apply mathematical transform
  - More complex distribution  $p(x)$ ?



# Sampling Methods

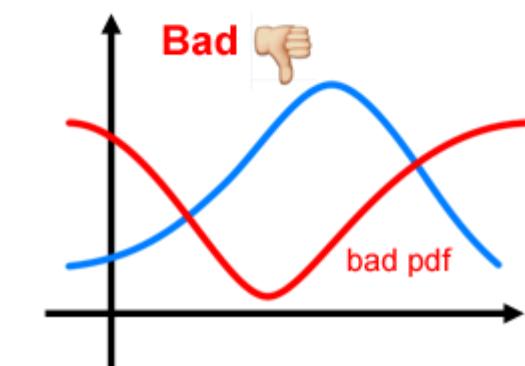
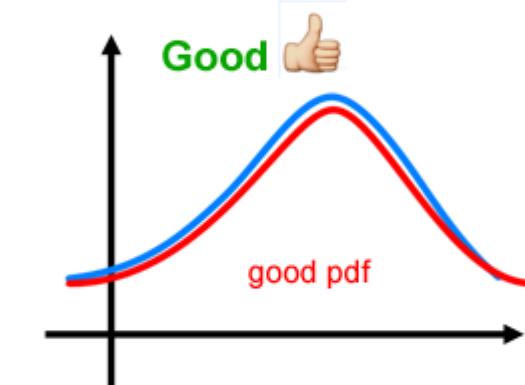
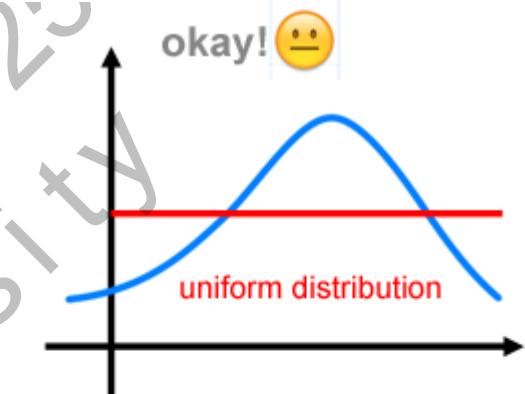
- Goal: sampling from  $p(x)$ 
  - No CDF or nice mathematical property available
- Idea: weighted samples
  - sample from “easy” distribution  $q(x)$  (e.g., uniform)
  - Use  $p(x)/q(x)$  as the weight for the sample
- Importance Sampling
  - $q(x)$  proposal distribution
  - $\frac{p(x)}{q(x)}$  importance weight
  - $E_{x \sim p}[f(x)] = E_{x \sim q} \left[ \frac{p(x)}{q(x)} f(x) \right]$



# Sampling Methods

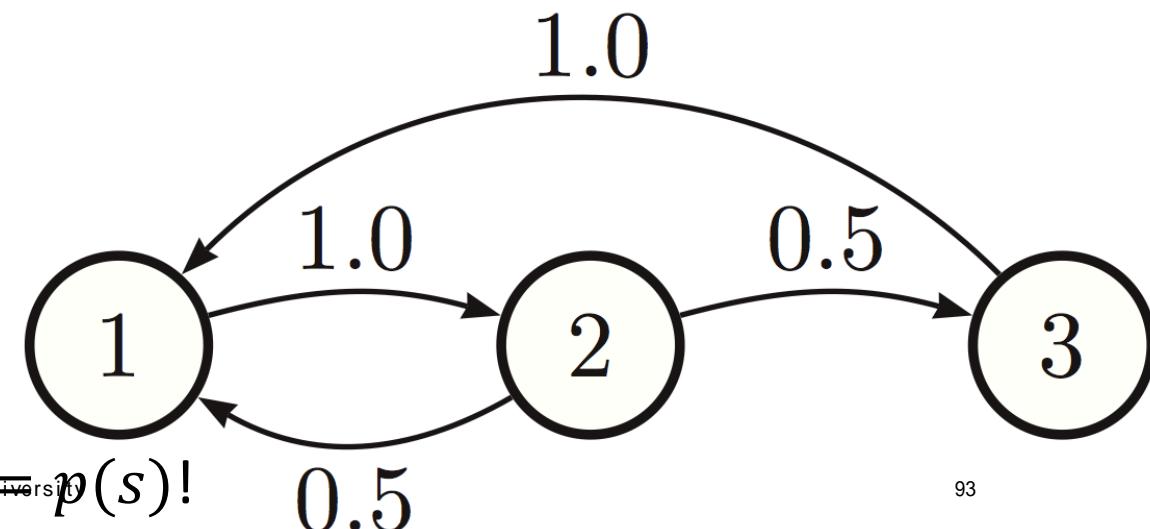
- Goal: sampling from  $p(x)$ 
  - No CDF or nice mathematical property available
- Idea: weighted samples
  - sample from “easy” distribution  $q(x)$  (e.g., uniform)
  - Use  $p(x)/q(x)$  as the weight
- Importance Sampling
  - $q(x)$  proposal distribution
  - **How to choose  $q(x)???$**
  - $q(x)$  needs to similar to  $p(x)$ 
    - Your homework ☺

**What if we don't have a universally good proposal?**



# Markov Chain Monte-Carlo

- Markov Chain
  - A state space  $S$ , a transition probability  $P(s_j | s_i) = T_{ij}$
  - $T$  is the transition matrix
  - We also use  $T(s_i \rightarrow s_j)$  to denote  $T_{ij}$
- A Markov Chain has a stationary distribution with a proper  $T$ 
  - Current distribution over states  $\pi_t$
  - Single step transition  $\pi_{t+1} = T\pi_t$
  - Stationary distribution  $\pi = T^\infty \pi_0$
- Sampling is easy!
- Goal: construct a Markov Chain
  - With a desired stationary distribution  $\pi = p(s)!$



# Markov Chain Monte-Carlo

- How to ensure  $\pi$  is a stationary distribution of a Markov Chain?
  - Detailed Balance (sufficient condition)
$$\pi(s)T(s \rightarrow s') = \pi(s')T(s' \rightarrow s)$$

# Markov Chain Monte-Carlo

- How to ensure  $\pi$  is a stationary distribution of a Markov Chain?
  - Detailed Balance (sufficient condition)
$$\pi(s)T(s \rightarrow s') = \pi(s')T(s' \rightarrow s)$$
  - **Design** a Markov chain satisfying detailed balance for desired density  $p(s)$ !

# Markov Chain Monte-Carlo

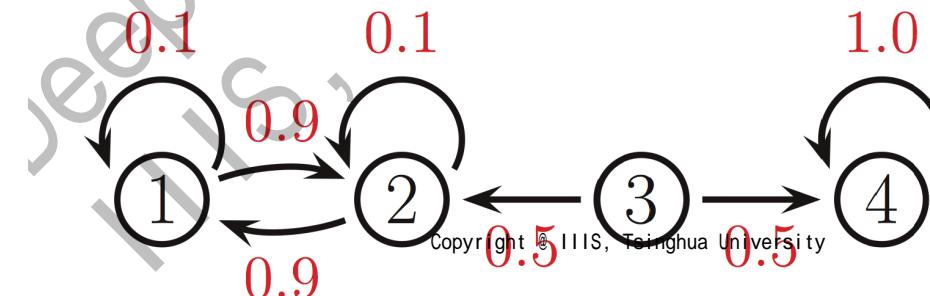
- How to ensure  $\pi$  is a stationary distribution of a Markov Chain?
  - Detailed Balance (sufficient condition)
 
$$\pi(s)T(s \rightarrow s') = \pi(s')T(s' \rightarrow s)$$
  - Design a Markov chain satisfying detailed balance for desired density  $p(s)$ !
- How to ensure a **unique** stationary distribution exist?

- The Markov chain is ergodic (遍历性) !

$$\min_z \min_{z':\pi(z')>0} \frac{T(z \rightarrow z')}{\pi(z')} = \delta > 0$$

*Intuitively: you can visit any desired state with positive probability from any state*

- Examples:



$$T = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

# Metropolis Hastings Algorithm

- Construct a valid Markov Chain  $T(s' \rightarrow s)$  for distribution  $p(s)$ 
  - Detailed balance:  $p(s)T(s \rightarrow s') = p(s')T(s' \rightarrow s)$
  - Ergodicity
- Metropolis Hastings Algorithm
  - A proposal distribution  $q(s'|s)$  to produce next state  $s'$  based on  $s$
  - Draw  $s' \sim q(s'|s)$
  - $\alpha = \min\left(1, \frac{p(s')q(s' \rightarrow s)}{p(s)q(s \rightarrow s')}\right)$  ( $q(s \rightarrow s')$  to denotes  $q(s'|s)$  for simplicity)
  - Transition to  $s'$  (**accept**) with probability  $\alpha$  (**acceptance ratio**);
  - O.w., stays at  $s$  (**reject**)
- MH constructs a valid Markov chain with a proper proposal  $q$ !
  - Homework ☺

# Metropolis Hastings Algorithm: Example

- Choice of  $q(s \rightarrow s')$ 
  - Random proposal  $q(s \rightarrow s') = s + \text{noise}$  (i.e., Gaussian/Uniform Noise)
- Acceptance ratio for  $s \rightarrow s'$ 
  - $\alpha(s \rightarrow s') = \min\left(1, \frac{p(s')q(s' \rightarrow s)}{p(s)q(s \rightarrow s')}\right) = \min\left(1, \frac{p(s')}{p(s)}\right)$
- MH sampling for the energy-based model  $p(s) = \frac{1}{Z} \exp(-E(s))$ 
  - Random initialize  $s^0$
  - $s' \leftarrow q(s \rightarrow s')$
  - Transition to  $s'$  with probability  $\alpha(s \rightarrow s')$ ;
  - O.w., stays at  $s$
  - Repeat

# Metropolis Hastings Algorithm: Example

- Choice of  $q(s \rightarrow s')$ 
  - Random proposal  $q(s \rightarrow s') = s + \text{noise}$  (i.e., Gaussian/Uniform Noise)
- Acceptance ratio for  $s \rightarrow s'$ 
  - $\alpha(s \rightarrow s') = \min\left(1, \frac{p(s')q(s' \rightarrow s)}{p(s)q(s \rightarrow s')}\right) = \min\left(1, \frac{p(s')}{p(s)}\right)$
- MH sampling for the energy-based model  $p(s) = \frac{1}{Z} \exp(-E(s))$ 
  - Random initialize  $s^0$
  - $s' \leftarrow s + \text{noise}$
  - Transition to  $s'$  with probability  $\alpha(s \rightarrow s')$ ;
  - O.w., stays at  $s$
  - Repeat

# Metropolis Hastings Algorithm: Example

- Choice of  $q(s \rightarrow s')$ 
  - Random proposal  $q(s \rightarrow s') = s + \text{noise}$  (i.e., Gaussian/Uniform Noise)
- Acceptance ratio for  $s \rightarrow s'$ 
  - $\alpha(s \rightarrow s') = \min\left(1, \frac{p(s')q(s' \rightarrow s)}{p(s)q(s \rightarrow s')}\right) = \min\left(1, \frac{p(s')}{p(s)}\right)$
- MH sampling for the energy-based model  $p(s) = \frac{1}{Z} \exp(-E(s))$ 
  - Random initialize  $s^0$
  - $s' \leftarrow s + \text{noise}$
  - Transition to  $s'$  with probability  $\min\left(1, \frac{p(s')}{p(s)}\right)$ ; **No partition function involved!**
  - O.w., stays at  $s$
  - Repeat

# Metropolis Hastings Algorithm: Example

- Choice of  $q(s \rightarrow s')$ 
  - Random proposal  $q(s \rightarrow s') = s + \text{noise}$  (i.e., Gaussian/Uniform Noise)
- Acceptance ratio for  $s \rightarrow s'$ 
  - $\alpha(s \rightarrow s') = \min\left(1, \frac{p(s')q(s' \rightarrow s)}{p(s)q(s \rightarrow s')}\right) = \min\left(1, \frac{p(s')}{p(s)}\right)$
- MH sampling for the energy-based model  $p(s) = \frac{1}{Z} \exp(-E(s))$ 
  - Random initialize  $s^0$
  - For each iteration  $t$ 
    - $s' \leftarrow s^t + \text{noise}$
    - If  $E(s') < E(s^t)$ ; then accept  $s^{t+1} \leftarrow s'$
    - Else accept  $s^{t+1} \leftarrow s'$  with probability  $\exp(E(s^t) - E(s'))$
  - Repeat

# Metropolis Hastings Algorithm

- The simplest way to construct a valid Markov chain
  - Flexible, simple and general
  - **Quiz: proposal  $q$  in MH v.s. Importance Sampling**
    - A:  $q(s'|s)$  v.s.  $q(s)$ ; in MH,  $q$  generates local samples; in IS,  $q$  outputs “blind” guesses
- Issues
  - Curse of dimensionality: samples a completely new state
  - Acceptance ratio: what if acceptance rate is low?

# Metropolis Hastings Algorithm

- The simplest way to construct a valid Markov chain
  - Flexible, simple and general
  - **Quiz: proposal  $q$  in MH v.s. Importance Sampling**
    - A:  $q(s'|s)$  v.s.  $q(s)$ ; in MH,  $q$  generates local samples; in IS,  $q$  outputs “blind” guesses
- Issues
  - Curse of dimensionality: samples a completely new state
  - **Acceptance ratio: what if acceptance rate is low?**
- Can we design a proposal distribution  $q(s \rightarrow s')$  such that it always gets accepted?

# Gibbs Sampling

- Gibbs sampling
  - $s = (s_0, s_1, \dots, s_N)$ , we construct a coordinate-wise  $q(s_i \rightarrow s'_i)$
  - $q(s_i \rightarrow s'_i) = p(s'_i | s_{j \neq i})$  (conditional distribution)
- Dimensionality
  - Sample a single coordinate per step.
- Gibbs sampling always get accepted!
  - Acceptance ratio is always 1,  $\alpha(s_i \rightarrow s'_i) = 1$  Prove it in your homework 😊
- Assumption
  - An easy to sample conditional distribution
    - Conjugate Prior and Exponential Family ([https://en.wikipedia.org/wiki/Conjugate\\_prior](https://en.wikipedia.org/wiki/Conjugate_prior))
  - What if no closed-form posterior?
    - Learn a neural proposal to approximate the true posterior! 😊  
(meta-learning MCMC proposals, Wang, Wu, et al NIPS2018)

# Sampling Methods

- What we have learned so far ...
  - Importance Sampling
    - Simplest solution by any proposal distribution
  - Metropolis-Hastings algorithm
    - Good local proposal → high acceptance ratio
  - Gibbs sampling
    - Posterior is easy-to-sample
    - The “default” method for machine learning among 2002~2012
- General Issues for MCMC methods
  - Slow convergence due to sampling (recap: SGD v.s. GD)
  - Can we use gradient information for MCMC?
    - Energy function is differentiable!

# Stochastic Gradient MCMC

- MCMC with Langevin dynamics
  - “Bayesian learning via stochastic gradient langevin dynamics”
    - ICML 2011, Max Welling & Yee Whye Teh (ICML 2021 test-of-time award)
  - Given  $N$  data  $X_1, \dots, X_N$ , define  $p(\theta \rightarrow \theta')$  by
$$\theta' \leftarrow \theta + \frac{\epsilon_t}{2} \left( \nabla_{\theta} \log p(\theta) + \sum_i \nabla_{\theta} \log P(x_i | \theta) \right) + N(0, \epsilon_t I)$$
  - Condition for a valid Markov Chain
    - $\sum_t \epsilon_t = \infty$  and  $\sum_t \epsilon_t^2 < \infty$
    - Interpretation
      - (stochastic) gradient descent first ( $\nabla_{\theta}$  is large); MCMC around local minimum ( $\nabla_{\theta} \approx 0$ )
    - No need of MH acceptance rule
- Additional Reading:
  - Hamiltonian Monte Carlo (SGD with momentum): <https://arxiv.org/pdf/1701.02434.pdf>
  - [https://arogozhnikov.github.io/2016/12/19/markov\\_chain\\_monte\\_carlo.html](https://arogozhnikov.github.io/2016/12/19/markov_chain_monte_carlo.html)

# Summary

- Hopfield Network
  - The first generative neural network
  - Undirected complete graph
- Boltzmann Machine
  - A probabilistic interpretation of Hopfield Network
  - The first deep generative model
- Energy-Based
  - Extremely flexible and powerful, designed to be multi-modal
  - Hard to sample and learn
  - **Sampling is the core challenge!!**

# What's Next: Non-Sampling Methods

- Approximate Bayesian Inference
  - Variational Inference (next lecture ☺)
    - Learn an parameterized distribution to approximate the true posterior
- Design a model from which we can easily draw sample!
  - Lectures 6 & 7a
- Modern energy-based models
  - Scoring matching
  - Lecture 7b



Song et. al., 2021  
OpenAI Blog: <https://openai.com/blog/energy-based-models/>

Thanks!

Deep Learning, Spring 2025  
IIIS, Tsinghua University